CMF 3D-QSAR

The continuous molecular fields approach, as a program

Desktop application, library and command-line tool
for the Continuous Molecular Fields approach to 3D-QSAR
User manual

`github.com/igbaskin/cmf-qsar`

Contents

What this program is 5

- The idea in one page 5
- What you can do with it 5
- How this manual is arranged 6

How a prediction is made 7

- 1. The fields 7
- 2. The overlap integral 7
- 3. Normalisation and scaling 8
- 4. Mixing the fields 8
- 5. Fitting: what the training set buys 8
- 6. Predicting a new molecule 9
- The same prediction, as a field 9
- A worked example 10
- What changes for the other tasks 10
- What has to match between fitting and predicting 11

Installation and first run 13

- Where the program lives 13
- Requirements 13
- Starting the application 13
- A first model in five minutes 14

The application window 15

- The Data set panel 15
- The Model panel 16
- The 3D view 17
- The Results panel 18
- The menus 19

Preparing the data 20

- What the program reads 20
- Activities 20
- Alignment 20
- Judging an alignment, and aligning flexibly 21
- Charges and per-atom properties 22
- Quantum-chemical properties with ORCA 22

Before you model: `cmf diagnose` 25

- Two things the command protects you from 25
- Reading the answer 26
- Does the advice work? 26

Building a model 27

Choosing the fields 27
GRID-like probe fields 27
Weighting the fields by solvent exposure 28
Descriptors that are not fields 29
Attenuation, regularisation and the grids 31
The learners 31
Fitting the order instead of the values: the ranking SVM 31
Nonlinear kernel layers and multiple kernel learning 33
Validation 33
Is it chance? y-randomisation 34
Where does the model apply? 34
Is the whole library inside? The two-sample test 35
How certain is that number? 36
How certain is *one* prediction? Conformal intervals 36
What the model thinks the groups are: clustering the kernel 37
Where the validation appears 38
Statistics reported 38
Saving and reusing 39

Reading the model 40

Co-fields 40
Contributions 40
Shapley values 40
A pharmacophore out of the co-field 41
What a candidate is missing 41
Two endpoints at once: the selectivity co-field 41
Chemical space maps 42

Classification and virtual screening 43

Setting up a classification task 43
One-class models: screening with actives alone 43
Decoys and the figures that matter 43

Conformations: the fourth dimension 45

Working without an alignment 46

Why the frame is the problem 46
The pose-free count kernel 46
The rotation-invariant autocorrelation kernel 46
Molecules on a surface: the in-plane invariant kernel 48
Crystals: the periodic autocorrelation kernel 50
Topomer fields 51
What the three routes are worth, measured 52
Heterogeneous series, step by step 52

Protein pockets 53

Getting a pocket on screen 53

Complementarity as numbers 53
Docking 54
Command-line reference 55
The options of `build`, by group 55
Worked examples 56
Using it as a library 57
The example data sets 58
When something goes wrong 59
Appendix A — Field parameterisation 60
Appendix B — Files and tags the program understands 61
Appendix C — Glossary 62
Appendix D — Bibliography 66
The method 66
One-class classification and virtual screening 66
Methods and data the program borrows 66

What this program is

CMF 3D-QSAR is a desktop application, a Python library and a command-line tool implementing the **Continuous Molecular Fields** approach to three-dimensional structure–activity modelling. The method was introduced by Zhokhova, Baskin and co-workers (Dokl. Chem. **429** (2009) 273; J. Comput.-Aided Mol. Des. **27** (2013) 427; continuous indicator fields: J. Comput.-Aided Mol. Des. **29** (2015) 233; co-fields: Future Med. Chem. **11** (2019) 2701), and the program follows those papers closely enough that their published statistics can be reproduced from the same data — `docs/benchmark.md` records that comparison set by set.

The idea in one page

A molecule is represented not by a table of descriptors and not by values on a lattice, but by **continuous functions of space**. Each field is a sum of Gaussians centred on the atoms,

$$X_f(\mathbf{r}) = \sum_l w_{lf} \exp\left(-\frac{\alpha_f}{2} |\mathbf{r} - \mathbf{r}_l|^2\right)$$

where w_{lf} is the weight of atom l in field f — its partial charge for the electrostatic field, its Crippen logP contribution for the hydrophobic field, 1 for an indicator field of a given atom type, or any number the file carries per atom — and α_f is the attenuation factor that sets how far the field reaches.

Two molecules are compared by the **overlap integral** of their fields, which for Gaussians has a closed form:

$$K_f(A,B) = \int X_{fA}(\mathbf{r})\, X_{fB}(\mathbf{r})\, \mathrm{d}\mathbf{r} = \left(\frac{\pi}{\alpha}\right)^{3/2} \sum_l \sum_m w_{lf} w_{mf} \exp\left(-\frac{\alpha}{4} |\mathbf{r}_l - \mathbf{r}_m|^2\right)$$

No grid is ever evaluated for this: the kernel is a sum over atom pairs, computed exactly. The kernels of the individual fields are mixed with coefficients h_f into one kernel, and that kernel goes into a kernel method — ridge regression, a support vector machine, a Gaussian process. Three consequences follow, and they are what distinguishes the approach from CoMFA:

- **Nothing is discretised.** There is no lattice spacing to choose, no alignment of grids, no variable selection among tens of thousands of columns, and no truncation at a cut-off distance.
- **The model is a function of space again.** A kernel ridge model predicts $y(X) = \Sigma_j a_j K(X, X_j) + b$, and because the kernel is an integral of a product, the whole sum can be rewritten as the overlap of the molecule's field with one function, the **co-field** $C_f(\mathbf{r}) = \Sigma_j a_j X_{fj}(\mathbf{r})$. The co-field is the program's map: it plays the role of CoMFA's coefficient contour, but it is the model rather than a projection of it.
- **The decomposition is exact.** Because the prediction is bilinear in the fields, it can be split without approximation into contributions of training molecules, of field types, of regions of space, or of individual atoms.

What you can do with it

Task	How
Build a 3D-QSAR regression model on an aligned congeneric series	Model panel, or `cmf_cli.py build`
Work without an alignment	topomer fields, the rotation-invariant autocorrelation kernel, or the pose-free count kernel
Classify actives against inactives, or screen with actives alone	binary (SVC, ν-SVC, Laplace GP) and one-class (1-SVM, centroid, k-NN, GP, kernel PCA) tasks
See what the model says about space	co-fields as isosurfaces, molecular fields, overlap maps, Gaussian-cube export
See what the model says about a molecule	training-molecule contributions, atom contributions, Shapley values and interaction indices
Turn a model into design instructions	pharmacophore extraction from co-field lobes, field-difference maps against a reference ligand
Use the protein	fetch or load a pocket, draw it behind the fields, score field complementarity, dock
Average over conformations	conformational ensembles (4D) and multi-instance learning
Map a data set	kernel GTM chemical-space maps with activity, class and density landscapes

How this manual is arranged

Chapter 2 follows one prediction from the atoms to the number, which is the method in full. Chapters 3 and 4 get you running and name every part of the window. Chapters 5 to 7 are the working cycle: prepare the data, fit a model, read it. Chapters 8 to 11 are the specialised routes — screening, conformations, alignment-free modelling, protein pockets. Chapter 12 is the complete command-line reference, chapter 13 the library, chapter 14 the example data, and chapter 15 collects the things that go wrong. The appendices hold the field parameterisation, the file conventions, a glossary of the terms this manual uses in a particular sense, and the bibliography.

Conventions: *Menu → Item* names a menu command, `monospace` names a file, an option or a piece of code. Equations are quoted in the notation of the papers, with their equation numbers.

How a prediction is made

This chapter follows one number — the activity predicted for one molecule — from the atoms it starts in to the value that appears in the table. Nothing here is a sketch: every formula is the one the program evaluates, and the worked example at the end was computed by running it.

The chain has six links:

1. each molecule becomes a set of **fields**, one per field type, each a sum of Gaussians on its atoms;
2. every pair of molecules gets a **field kernel** per field type, the overlap integral of those two fields, computed in closed form;
3. each field kernel is **normalised** and **scaled**;
4. the field kernels are **mixed** into one kernel matrix;
5. the learner turns that matrix and the activities into **dual coefficients**;
6. a new molecule is pushed through steps 1 to 4 against the training set, and its prediction is the **weighted sum of its kernel values** with those coefficients.

1. The fields

Atom l of a molecule carries a weight w_{lf} in field f — its partial charge in the electrostatic field, its Crippen logP contribution in the hydrophobic field, an Abraham-type increment in the hydrogen-bond fields, the Tripos van der Waals well depth in the steric field, 1 in an indicator field for atoms of that type and 0 elsewhere, or a number read from the file in a `prop:` field. Appendix A lists them.

The field itself is the sum of Gaussians those weights define,

$$X_f(\mathbf{r}) = \sum_l w_{lf} \exp\left(-\frac{b_{lf}}{2}|\mathbf{r}-\mathbf{r}_l|^2\right)$$

with width $b_{lf} = \alpha_f$ for every field except the steric one, where it is scaled by the atom's van der Waals radius, $b_{lf} = \alpha_f / R_l^2$, so that a large atom spreads its bulk over a correspondingly larger region.

There is no grid. The field is never evaluated anywhere in order to make a prediction; it is evaluated only when a map has to be *drawn*.

2. The overlap integral

Two Gaussians have an overlap integral in closed form, so the inner product of two whole fields is a double sum over atom pairs:

$$K_f(A,B) = \int X_{fA}(\mathbf{r}) X_{fB}(\mathbf{r})\,\mathrm{d}\mathbf{r} = \sum_{l\in A}\sum_{m\in B} w_{lf} w_{mf} \left(\frac{2\pi}{b_{lf}+b_{mf}}\right)^{3/2} \exp\left(-\frac{b_{lf}b_{mf}}{2(b_{lf}+b_{mf})}|\mathbf{r}_l-\mathbf{r}_m|^2\right)$$

When every atom shares one width, $b_{lf} = b_{mf} = \alpha$, this collapses to the form the papers quote,

$$K_f(A,B) = \left(\frac{\pi}{\alpha}\right)^{3/2} \sum_l \sum_m w_{lf} w_{mf} \exp\left(-\frac{\alpha}{4}|\mathbf{r}_l-\mathbf{r}_m|^2\right)$$

and the general form is what the steric field needs, since its widths differ atom by atom. A unit test checks both against numerical quadrature of the fields on a fine grid.

Two things follow immediately. The cost of a kernel is the number of atom pairs, not the number of grid points, which is why a 300-molecule set takes seconds. And the attenuation α is the only thing that decides how far apart two atoms may be and still contribute: large α demands close matching, small α lets distant atoms see each other, and in the limit α → 0 the geometry drops out altogether and the kernel becomes a product of atom counts — which is the alignment-free baseline of *Working without an alignment*.

3. Normalisation and scaling

Two corrections stand between the raw integrals and a usable kernel, and they do different jobs.

Cosine normalisation (*Normalise kernels*) divides each field kernel by the norms of the two fields,

$$\hat{K}_f(A,B) = \frac{K_f(A,B)}{\sqrt{K_f(A,A)\,K_f(B,B)}}$$

so that every molecule has unit self-similarity in every field -- the normalised kernel of a molecule with itself is 1 by construction -- and a large molecule cannot look important merely for being large.

Field scaling (*Scale fields*) then multiplies each field kernel by a constant c_f chosen so that its mean diagonal over the training set is 1. It equalises the fields' contributions before they are mixed, which matters when one field's raw magnitudes dwarf another's — an electrostatic kernel built from partial charges and a steric kernel built from well depths are not naturally on the same scale. With cosine normalisation on, every diagonal is already 1 and c_f = 1; the two options are therefore alternatives in practice, and unticking both gives the raw-kernel convention of the papers.

The scaling constants are computed **on the training set and then stored**: a prediction must not rescale anything according to the molecules being predicted, or two test molecules would be scored on different scales.

4. Mixing the fields

The joint kernel is a weighted sum over field types,

$$K(A,B) = \sum_f h_f c_f \hat{K}_f(A,B)$$

with the mixing coefficients h_f left at 1 by default, edited by hand in the field table, chosen by coordinate descent on the cross-validated error, or set by centred kernel alignment. Since each normalised field kernel has unit diagonal, the joint self-kernel of any molecule is $\Sigma_f\, h_f\, c_f$ — five, for the five default fields with unit coefficients, which is the 5 on the diagonal of the matrix in the worked example below.

5. Fitting: what the training set buys

Kernel ridge regression with a bias term solves one linear system (Eq. 12 of the 2013 paper):

$$(\mathbf{K} + \lambda \mathbf{I})\,\mathbf{a} + b\,\mathbf{1} = \mathbf{y}, \qquad \mathbf{1}^{\mathrm{T}}\mathbf{a} = 0$$

The outcome is one **dual coefficient** a_j per training molecule and one bias b. The last row of the system is the constraint $\Sigma_j\ a_j = 0$, which is what makes the bias meaningful: the coefficients describe only *departures* from a constant level, and b carries the level itself. The ridge parameter λ adds to the diagonal, which keeps the system solvable when molecules are nearly identical — as they are in a congeneric series, where the off-diagonal kernel values run above 0.97 — and controls how hard the model chases the training activities.

The other learners differ only in how **a** and b are obtained, never in what they are: ε-SVR and ν-SVR get them from a quadratic program in which most a_j come out exactly zero, C-SVC and ν-SVC the same for class labels, and the Laplace Gaussian-process classifier from the posterior mode of a latent function. In every case the fitted model is the pair (**a**, b) over the training molecules.

6. Predicting a new molecule

For a molecule the model has never seen, the program:

1. builds its fields **with the same field set** — the same α, the same charge model, the same parameterisation, the same atom typing;
2. computes the per-field overlap integrals between its atoms and the atoms of every training molecule, giving one row $K_f(x, j)$ per field;
3. normalises that row with the molecule's own self-kernels and the **stored** training self-kernels, and multiplies by the **stored** scales c_f;
4. mixes those normalised, scaled fields into one row $k(x)_j$ with the mixing coefficients h_f, which is the joint kernel $K(x, j)$ of *Mixing the fields* above;
5. returns

$$\hat{y}(x) = \sum_j a_j\, k(x)_j + b = \mathbf{k}(x)^{\mathrm{T}}\mathbf{a} + b$$

That is the whole prediction: a similarity-weighted combination of the training activities, where "similarity" is the overlap of molecular fields and the weights were fixed when the model was fitted. Molecules with $a_j = 0$ — most of them, after an SVM fit — contribute nothing, which is why the *Contributions* tab can list the handful that actually make a given prediction.

The same prediction, as a field

Because the kernel is an integral of a product and the prediction is linear in it, the sum over training molecules can be moved inside the integral. Writing $s_j = 1/\sqrt{(K_f(j,j))}$ for the normalisation of training molecule j, define the **co-field**

$$C_f(\mathbf{r}) = \sum_j a_j\, s_j X_{fj}(\mathbf{r})$$

and the prediction becomes

$$\hat{y}(x) = \sum_f \frac{h_f c_f}{\sqrt{K_f(x,x)}} \int X_{fx}(\mathbf{r})\, C_f(\mathbf{r})\, \mathrm{d}\mathbf{r} + b$$

The model is therefore a *function of space*, one per field type, and a molecule is scored by how much of its own field lies in the positive regions of that function and how much in the negative

ones. This is the identity that makes the pictures of *Reading the model* the model rather than an illustration of it, and the division by the molecule's own field norm is the only thing the normalisation adds to the picture. It also explains what co-fields cannot survive: a nonlinear kernel layer, which is not an integral of a product, and the autocorrelation kernel source, which has no frame to draw in.

A worked example

Four molecules from `examples/benzamides_synthetic.sdf` as the training set, the five default fields, α = 0.5, cosine normalisation, λ = 0.1, and the fifth molecule as the one to predict. The numbers below are the program's own, printed from the objects it fits.

The joint kernel of the training set, the sum of the five normalised field kernels, so the diagonal is 5:

```
                parent     4F      4'F     4Cl          y
parent          5.0000  2.8362  3.2195  2.5368       5.585
4F              2.8362  5.0000  2.3310  3.9647       5.230
4'F             3.2195  2.3310  5.0000  2.8882       6.471
4Cl             2.5368  3.9647  2.8882  5.0000       5.355
```

Solving the bordered system gives

```
a = [-0.1358, -0.0581, +0.3487, -0.1549]      sum(a) = 0      b = 5.7124
```

The test molecule 4'Cl is then compared with each training molecule, field by field:

```
                        parent      4F      4'F     4Cl
electrostatic           0.0878  0.6028  0.1009  0.3915
steric                  0.9793  0.9824  0.9863  0.9579
hydrophobic             0.8936  0.9412  0.9443  0.8437
H-bond acceptor         0.2600  0.9895  0.1927  0.6539
H-bond donor            0.5815  1.0000  0.4424  0.7379
                        ------  ------  ------  ------
k(x)  (their sum)       2.8022  4.5159  2.6667  3.5849
```

and the prediction is that row times the coefficients, plus the bias:

```
y(x) = 2.8022(-0.1358) + 4.5159(-0.0581) + 2.6667(0.3487) + 3.5849(-0.1549) + 5.7124
     = 5.4444
```

which is exactly what *Pred (test)* shows for that molecule. Two things are worth reading off this example rather than taking on trust. The steric and hydrophobic rows are nearly constant across the training set — these molecules differ by one halogen — so almost all the discrimination comes from the electrostatic and hydrogen-bond fields, which is visible in the spread of their rows. And the prediction, 5.44 against an observed 7.21, is poor, because four training molecules and no hyper-parameter search is not a model: the arithmetic is the lesson here, not the accuracy.

What changes for the other tasks

Model	What the prediction is
Kernel ridge	$\hat{y} = \mathbf{k}^{\mathsf{T}}\mathbf{a} + b$, as above

Model	What the prediction is
Gaussian process	the same mean, plus the variance $\sigma^2(x) = K(x,x) - \mathbf{k}^{\mathsf{T}}(\mathbf{K}+\lambda\mathbf{I})^{-1}\mathbf{k}$, which the *Align by prediction confidence* tool minimises
ε-SVR, ν-SVR	the same formula; most a_j are zero
C-SVC, ν-SVC	the same formula, read as a decision value: its sign is the class, its magnitude the margin
Laplace GP classifier	the same formula gives the latent mean, which is squashed to a probability and averaged over the latent variance
One-class methods	a score built from the same kernel row, but also from the molecule's own self-kernel: distance to the centroid, to the k-th nearest active, to the principal subspace, or the 1-SVM decision value
Nonlinear kernel layer	the joint kernel is passed through $\exp(-\gamma\, d^2)$ or an arc-cosine function first; the prediction formula is unchanged, the co-field identity is lost
Autocorrelation source	the kernel row comes from the rotation-invariant distance profiles instead of the overlap integrals; everything downstream is identical
Surface source	for molecules over a surface at $z = 0$: the in-plane distance profiles resolved by the heights of both atoms, invariant in the plane and to rotation about the normal only; heights and tilts are kept, no in-plane alignment is needed (command line: `--kernel surface --surface-z ZMIN ZMAX ZSTEP`)
Periodic source	for crystals: the distance profiles summed over the neighbour list of the cell with its images, per atom, background removed, cutoff tapered; invariant to the choice of cell, origin and orientation (command line: `build-crystals`)
Two endpoints	the same formula once per endpoint, with coefficients tied through a task correlation; the difference of the two coefficient vectors gives the selectivity co-field

What has to match between fitting and predicting

Because the kernel row is computed against the stored training atoms with the stored constants, a prediction is only meaningful if the new molecule was prepared the way the training molecules were. The saved `.cmf` model therefore carries the field set, the α values, the charge model, the parameterisation, the normalisation and scaling flags, the mixing coefficients, the training atoms and the coefficients — everything in steps 1 to 5 — and reuses them. It also carries a record of how the training molecules were prepared: the version of the program that wrote it, the date, how many molecules, where their partial charges came from (the file, or computed), whether they carried explicit hydrogens, and the box their atoms occupy. Before any prediction the program compares the new molecules with that record and reports what does not match:

- **The chemistry behind the weights.** Charges read from a file for the training set and Gasteiger charges for the new molecules are two different electrostatic fields wearing one name. A model fitted on file charges therefore **refuses** a molecule that carries none, rather than falling back to Gasteiger charges silently; `cmf_cli.py predict --allow-mismatch` overrides the refusal, and the application asks. Supply the charges (a MOL2 charge column or a `partial_charge` atom property) or build the model with computed charges.
- **Explicit hydrogens.** The donor field sits on the hydrogens and the steric field counts them, so a set with hydrogens and a set without are two different sets of fields. A mismatch is a **warning**.

- **The frame.** The overlap integral is not invariant to rotation or translation, so a molecule predicted in a different frame from the training set scores low on every field and is predicted towards the mean. A molecule whose centroid lies more than 4 Å outside the box of the training atoms is reported as a **warning**.
- **The pose inside the frame.** A molecule can be in the right place and still be turned the wrong way, and a rotation about its own centre moves no centroid, so the box test above sees nothing. The check that does see it compares the molecule with the *shape* the training set occupies rather than with its box: the share of its atoms that come within 1.5 Å of some training atom, against the share this training set asks of its own molecules — the fifth percentile of the same quantity measured on each training molecule against the others, less a small slack, so that a tight series earns a tight rule exactly as the applicability domain's threshold does. Measured on the eight Sutherland sets, turning every held-out molecule 90° about its own centre is caught **95 % of the time, for a false-positive rate of 2.6 %** on the molecules left as they were. If the new molecules are not in the frame, align them (*Preparing the data*) or use a frame-free kernel source (*Working without an alignment*).
- **Per-atom property fields** that are zero on a molecule, because it carries no value of that property, are a warning; a file written by another version of the program, or by a version that recorded no provenance, is a note.

The findings appear in the log of the application and on the console of the command-line tool, each marked `[refuse]`, `[warn]` or `[note]`. A model file written by a newer version of the program than the one reading it is not opened.

Installation and first run

Where the program lives

The repository is github.com/igbaskin/cmf-qsar — private, so the link asks for a GitHub account with access:

```
git clone https://github.com/igbaskin/cmf-qsar.git
cd cmf-qsar
```

Everything this manual describes is in that checkout: the `cmfqsar` package with the command-line tool inside it (`cmfqsar/cli.py`), the two launchers (`cmf_app.py`, `cmf_cli.py`), the example data under `examples/`, the benchmark scripts under `benchmarks/` with their recorded results, the tests, and this manual with the script that typesets it.

Requirements

Python 3.10 or newer. The package is declared in `pyproject.toml`, so from the checkout

```
pip install -e .[gui]
```

installs the library with RDKit (chemistry), scikit-image (marching cubes for the isosurfaces), scikit-learn (the SVM solvers), NumPy, SciPy and joblib, and with the `gui` extra PySide6 (the interface) and pyqtgraph with PyOpenGL (the 3D view); it also puts the commands `cmf-app` and `cmf` on the path. The library alone, `pip install -e .`, needs no Qt. The program is pure Python and needs no compilation step; it has been used on Windows, Linux and macOS.

Further extras, only needed for particular work:

- `docs` — reportlab and matplotlib (which supplies the DejaVu fonts) to typeset this manual;
- `dev` — pytest, for the test suite;
- `examples` — `dimorphite_dl`, protonation states when importing SMILES for a given pH;
- a local ORCA or VASP installation if you intend to compute per-atom quantum properties to use as fields (see *Preparing the data* and appendix A); these are programs, not packages.

`pip install -r requirements.txt` installs the working checkout, `.[gui,docs,dev]`. `requirements-lock.txt` pins the exact versions the benchmark tables were produced with: `pip install -r requirements-lock.txt` and then `pip install --no-deps -e .` reproduces that environment. `CITATION.cff` says how to cite the software and the papers.

Starting the application

```
cmf-app
```

or `python cmf_app.py` from the checkout, which takes the same arguments; straight onto a data set, which saves two dialogs:

```
cmf-app examples/benzamides_synthetic.sdf --activity pActivity --test-fraction 0.25
```

The command-line tool is `cmf` once installed and `python cmf_cli.py` from the checkout; the rest of this manual writes the latter, and the two are the same program.

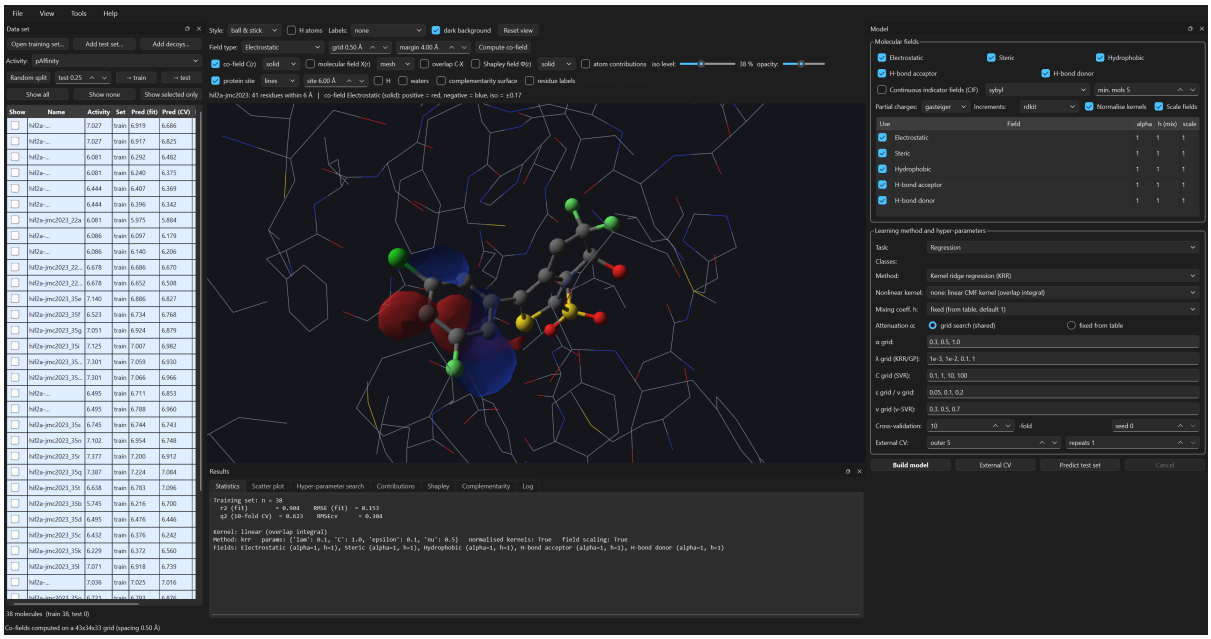
If the 3D view stays black, the OpenGL driver is the usual cause; see *When something goes wrong*.

A first model in five minutes

1. *File → Open training set…* and choose `examples/benzamides_synthetic.sdf`. It holds 109 aligned N-phenylbenzamides with a synthetic activity, so you know the answer the model should find. Pick `pActivity` as the activity property.
2. In the data-set table, press *Random split* and set a quarter of the molecules aside as a test set.
3. In the Model panel leave the five physico-chemical fields ticked, leave the method at kernel ridge regression, and press **Build model**.
4. Read the *Statistics* tab: q^2 should be about 0.87 and R^2_{pred} about 0.84 — the noise ceiling of that set is 0.94, so the model has recovered nearly all of the signal that was put into it.
5. Press **Compute co-field**, choose *Electrostatic*, and move the *iso level* slider. The red and blue lobes that appear are the pseudo-receptor zones the activity was generated from.

Nothing in that sequence is specific to the example: it is the working cycle of the program.

The application window



The HIF-2α series of the Merck free-energy benchmark, 38 ligands: a kernel ridge model with $q^2 = 0.62$, its electrostatic co-field drawn as ± isosurfaces, and the binding site of the reference crystal structure, which the data set brings with it and which needs no superposition.

The window is a 3D view with three dockable panels around it. The **3D view** and its control strip occupy the centre and are always there; the **Data set** panel docks on the left, the **Model** panel on the right inside a scroll area, and the **Results** panel across the bottom under the view. All three are ordinary dock widgets: drag one by its title bar to another edge, drop it on another to stack them as tabs, or pull it off entirely to float it on a second monitor. The layout is remembered between sessions, and *View → Restore default layout* puts it back when an experiment goes wrong.

At the bottom, the **status bar** shows what the program is doing, with a progress bar on the right during a long task. Every long task — a grid search, a co-field, an alignment, a conformer generation — runs in a worker thread, so the window stays responsive and the *Cancel* button in the model panel stops the work rather than the application.

The Data set panel

Three buttons at the top load structures: **Open training set…**, **Add test set…** and **Add decoys…**, the same actions as the first three entries of the *File* menu. Beside them the *Activity* selector names the SD property the activities are read from; changing it re-reads the column without reloading the file.

The table has one row per molecule and seven columns:

Column	Contents
Show	the tick that decides whether this molecule is drawn in the 3D view
Name	the molecule's title from the file
Activity	the experimental value, from the chosen property

Column	Contents
Set	train, test or decoy
Pred (fit)	the model's prediction for the molecule it was fitted on
Pred (CV)	its prediction when it was in the held-out fold
Pred (test)	its prediction as a test molecule
AD	its similarity to the training set, when the applicability domain was computed

An eighth column, **Chemotype**, appears once the chemotypes have been found. Rows are selected singly or in blocks, and the selection drives everything else: the co-field overlap, the atom contributions, the Shapley decomposition and the complementarity numbers all refer to the molecule currently selected.

Sorting. Clicking a column header sorts by it, a second click reverses the order, and a third restores the order of the file. Molecules with no value in that column always sink to the bottom, whichever direction is in force, so an unpredicted molecule never hides among the predictions. The *Show* ticks, the selection and the 3D view follow the molecule and not the row, so sorting never disturbs what is on screen.

Splitting. *Random split* assigns a fraction to the test set; the → *train* and → *test* buttons move the selected rows by hand. A `subset`, `set` or `split` tag in the file is honoured on loading, which is how the published splits of the example sets arrive without any clicking.

Showing. *Show all*, *Show none* and *Show selected only* control the view in bulk. With one molecule shown, the maps in the viewer refer to it; with all of them shown, the co-field sits in the middle of the series as it should.

A line under the table reports the state of the data set — how many molecules, how many in each subset, which activity property — and a note in brackets records any transformation that has been applied, for instance *[topomer substituents of …]*, so that a changed data set never passes for the original.

The Model panel

Everything that defines a model, in the order one normally works down it. The panel scrolls; on a small screen nothing is hidden, it is merely below.

Molecular fields. Five tick boxes for the physico-chemical fields — Electrostatic, Steric, Hydrophobic, H-bond acceptor, H-bond donor — then *Continuous indicator fields (CIF)* with its typing scheme (`sybyl`, `element`, `element+aromatic`) and the minimum number of training molecules an atom type must occur in to earn a field. Below them, *Per-atom property fields* lists the per-atom properties the loaded file carries, each of which can be turned into a field of its own.

Partial charges chooses between Gasteiger, MMFF94 and the charges stored in the file; *Increments* chooses the parameterisation (`rdkit`, `tripos`, `tripos-legacy`). Two tick boxes set the kernel conventions: *Normalise kernels* for cosine-normalised field kernels, *Scale fields* to rescale every field kernel to unit mean diagonal before mixing — untick it for the raw-kernel convention of the papers.

Conformations switches between one conformation per molecule and an ensemble, with the weighting (uniform or Boltzmann), kT, a cap on the number of conformations, and the multi-instance settings (the softness β, and *one conformation* for the hard selection). A label beside

them reports how many conformations the data set currently holds.

A table at the foot of the group lists the fields that will actually be built, with the columns **Use**, **Field**, **alpha**, **h (mix)** and **scale**: per-field attenuation and mixing coefficients are editable there, which is how a field is given more or less weight than the others without leaving the panel.

Learning method and hyper-parameters. A form, row by row:

Row	What it sets
Task	regression, binary classification, or one-class (virtual screening)
Classes	for classification: a threshold on the activity, with *active below* for data where small is potent, or a class property from the file
Method	the learner; the list changes with the task — KRR, GP, ε-SVR, ν-SVR for regression; C-SVC, ν-SVC, Laplace GP for classification; 1-SVM, centroid, k-NN, one-class GP and kernel PCA for one-class work
Nonlinear kernel	none (the linear CMF kernel), a Gaussian layer, or one or two arc-cosine layers
Mixing coeff. h	fixed from the table, optimised by coordinate descent, or set by centred kernel alignment
Attenuation α	a single value, or the grid below
Kernel from	overlap integrals of the fields, which need one frame; the rotation-invariant autocorrelation, which does not; or the surface autocorrelation for poses over a surface at $z = 0$, which needs only the normal (command line)
α grid, λ grid, C grid, ε/ν grids	the search ranges, as comma-separated lists
Cross-validation	the number of folds and the shuffling seed
External CV	the outer folds and the number of repeats for the nested procedure
Probe fields	GRID-like fields for the water, methyl, carbonyl, amine and cation probes, and the number of Gaussians per atom
Computed descriptors	molecular weight, clogP, TPSA and rotatable bonds, computed by RDKit; all off by default
Descriptors (per molecule)	which numeric molecule properties enter as non-field descriptors, and the block's mixing coefficient
Validation	*shuffles* for y-randomisation, *bootstrap* for confidence intervals, and the *applicability domain* tick — all three off by default, all three run as part of **Build model** rather than from a button of their own

Four buttons close the panel: **Build model** runs the grid search, the final fit, the internal cross-validation and the test prediction; **External CV** runs the nested procedure and reports q^2_{ex}; **Predict test set** applies the current model to the test molecules; **Cancel** stops whatever is running.

The 3D view

The view itself is an OpenGL scene: drag to orbit, the wheel to zoom, middle-drag (or shift-drag) to pan. *Reset view* reframes the whole scene, *Tools → Center view on current molecule* reframes the selected one, and *File → Save 3D view as image…* writes what you see to PNG.

The control strip above the view is four rows, and everything on it redraws immediately:

- **Molecules.** Drawing style (ball & stick, sticks, lines, space fill), *H atoms*, *Labels* (none, element + number, atom type, partial charge), *dark background*, *Reset view*.
- **Maps.** *Field type* chooses which field the maps refer to; **Compute co-field** evaluates the co-fields on a grid around the training set; then the maps themselves — *co-field C(r)*, *molecular field X(r)*, *overlap C·X*, *Shapley field Φ(r)*, *atom contributions*, *difference map*, *pharmacophore* — each with its own solid-or-mesh style where that applies. After a two-endpoint fit an *Endpoint* box joins the row and chooses whose co-field the maps show: either endpoint, or their difference, the selectivity co-field.
- **Levels.** *iso level*, the isosurface threshold as a percentage of the largest absolute value rather than an absolute number, so the surfaces stay comparable when the field changes; and *opacity*.
- **Protein site.** *protein site* draws the residues within the chosen distance of the molecules on screen, as lines or sticks, with *H*, *waters* and residue labels optional, and *complementarity surface* paints the molecule's surface by how well it complements the pocket.

Co-field colours are fixed, so that maps made in different sessions can be compared at a glance:

Field	Positive lobe	Negative lobe
Electrostatic	red	blue
Steric	green	yellow
Hydrophobic	orange	light blue
H-bond acceptor	magenta	dark red
H-bond donor	cyan	purple
Indicator field	red: the atom type is welcome here	blue: it is not

The Results panel

Twelve tabs, each filled by the work that produces it:

Tab	What it holds
Statistics	every figure of merit, for the fit, the cross-validation, the test set and the external CV; the y-randomisation, bootstrap and applicability-domain results follow at the bottom, under a *Validation* heading, when they were asked for
Scatter plot	observed against predicted, by subset, with the molecule under the cursor named
ROC	the receiver-operating curves of a classification model, per subset
Scores	the ranked decision values or calibrated probabilities: rank, molecule, score, set, class
Enrichment	the enrichment curve against a decoy set, with EF at 0.5, 1, 5 and 10 %, BEDROC and RIE
Map	the kernel-GTM chemical-space map, with the activity, class and density landscapes, and a control for how the molecules themselves are drawn over it
Hyper-parameter search	the whole grid — α, the other parameters, RMSEcv and q^2 — so that an optimum at the edge of the grid is visible rather than hidden

Tab	What it holds
Contributions	two tables for the selected molecule: the training molecules that make its prediction, with `a_j` and `a_j K(X, X_j)`, and its own atoms with their contributions
Shapley	the Shapley decomposition, over atoms, substituents of the common scaffold or field types, with second-order interactions
Complementarity	the field complementarity of the molecules and of the co-field with the pocket, by sign agreement, cosine or correlation
Chemotypes	the Murcko scaffolds found in the set, their sizes, and the leave-one-chemotype-out result beside the alignment-free baseline
Log	everything the program did, with timings and warnings — including the names of the molecules that fell outside the applicability domain — the first place to look when a result surprises you

The menus

File — *Open training set…* (Ctrl+O), *Add test set…*, *Add decoys (background set)…*, *Import activities from CSV…*; *Load protein structure (PDB/mmCIF)…*, *Fetch protein from the PDB…*, *Forget protein structure*; *Load model…* (Ctrl+L), *Save model…* (Ctrl+S); *Export predictions (CSV)…*, *Export aligned data set (SDF)…*, *Export current co-field (Gaussian cube)…*, *Export current Shapley field (Gaussian cube)…*, *Save 3D view as image…*; *Quit* (Ctrl+Q).

View — show or hide each panel, *Results panel spans full width* for when the tables and plots matter more than the view, and *Restore default layout*.

Tools — *Generate conformations (4D)…*; *Dock into the pocket…*; *Align data set to reference…*; *Align by prediction confidence (GP variance)…* and its cross-validated form for the training set; *Replace molecules by their topomer substituents…*; *Field difference from a reference molecule…*; *Extract pharmacophore from the co-field…*; *Model two endpoints (selectivity co-field)…*; *Chemical space map…*; *Align decoys to the training set…*; *Unalign (randomise poses)…*; *Center view on current molecule*; *Quantum-chemical properties with ORCA…* and *Save the data set with its properties (SDF)…*; and the submenu *Heterogeneous sets (several chemotypes)* with its four numbered steps.

Help — *About CMF 3D-QSAR / method summary*, which names the equations of the papers that each part of the program implements.

Every *Tools* entry opens a dialog that states what it is about to do and on which molecules, and each one runs in the background with the progress bar and the *Cancel* button live.

Preparing the data

What the program reads

Format	Notes
`.sdf`, `.sd`, `.mol`	3D coordinates used as they are; hydrogens are added with coordinates if missing
`.mol2`	Sybyl atom types and partial charges are read and can be used directly; a tolerant fallback reader handles blocks RDKit refuses
`.smi`, `.csv`, `.tsv`, `.txt`	2D input: each molecule is embedded with ETKDG and minimised with MMFF94, which means **the poses are arbitrary** until you align them
`.pdb`, `.mmcif`	protein structures only, read by a deliberately permissive reader

A CSV needs a column of SMILES; a column called `name` or the first non-SMILES column names the molecules, and every other column is kept as a molecule property, so activities and metadata travel with the file.

Activities

Three ways, in order of convenience:

1. **An SD property.** The loader guesses the activity property and the dialog lets you change it. Anything numeric will do: `pIC50`, `pKi`, `pActivity`, `CBG_pK`.
2. **A separate table** (*File → Import activities from CSV…*, `--activity-file`): molecule name, activity, and optionally a third column with the split (1 = train, 2 = test, or the words).
3. **A class property** for classification tasks, or a threshold on the activity.

A per-molecule tag named `subset`, `set` or `split` is honoured automatically, which is how the published splits of the example sets arrive.

Alignment

CMF compares fields in one frame of reference. If your molecules came from a superposition, from a crystallographic series or from a docking run, that frame already exists and you should leave it alone. If they came from SMILES, they have no frame at all and you must make one — or use one of the frame-free routes of *Working without an alignment*, which is often the better answer.

The tools, in increasing order of sophistication:

- *Tools → Align data set to reference…* — rigid superposition on the maximum common substructure with a reference molecule, optionally refined by maximising the CMF kernel similarity, which is the alignment criterion proposed in the 2013 paper.
- *Tools → Align by prediction confidence…* — minimises the Gaussian-process posterior variance $\sigma^2(x) = K(x,x) - \mathbf{k}(x)^{\mathrm{T}}(\mathbf{K} + \lambda\mathbf{I})^{-1}\mathbf{k}(x)$ of a molecule under an existing model, which moves it into the model's applicability domain. It uses no activity value, so it is valid for prospective prediction. The program refuses to move training molecules, because the optimum there is degenerate: superimposing everything maximises mutual similarity and destroys the signal.

- *Tools → Align decoys to the training set…* — the same idea for a background set, so that no decoy is rejected merely for having been posed badly.
- *Tools → Unalign (randomise poses)…* — the control experiment. It gives every molecule a random rigid pose, which leaves internal geometry and self-kernels untouched and scrambles only the mutual arrangement. Randomise, realign, and see how much comes back; the program remembers the poses from before and reports the recovery.

Judging an alignment, and aligning flexibly

Everything in this manual says the alignment decides more than the kernel, the learner or the fields. What it could not say until recently is whether one alignment method is better than another, because the obvious measure cannot tell you: realigning changes the geometry q^2 is computed on, so a method is rewarded for making the data look like the model, and q^2 and the external R^2pred have been seen to move in opposite directions under exactly that operation.

The judge has to come from outside the model, and the FEP benchmark sets supply one: every ligand of `examples/openff` is co-crystallised in the same frame, so the pose a method *should* have found is known. `benchmarks/run_poses.py` fits a model on training molecules in their crystal poses, destroys the poses of the held-out ones, and asks each method to recover them without ever seeing a held-out activity. Two numbers come back for every method — how near the crystal pose it got, measured over heavy atoms, symmetry-aware and **without superposition**, and what the model then predicts — and both are reported between two fixed references: the crystal poses as the ceiling and the scrambled poses as the floor.

That last detail is the whole protocol. Superposing before measuring would report 0.00 Å for a molecule placed in the wrong site: on a ligand displaced by 1 Å, `CalcRMS` returns 1.0 and `GetBestRMS` returns 0.0.

What the methods are worth

Measured this way over sixteen FEP sets, with the held-out **conformations regenerated** rather than read from the crystal — which is the case a new molecule actually presents:

method	median RMSD	within 2 Å	R^2pred	conformer limit
crystal poses (the ceiling)	0.00 Å	100 %	+0.334	0.00 Å
scrambled (the floor)	6.12 Å	1 %	−0.696	1.73 Å
substructure (MCS)	1.94 Å	54 %	−0.391	1.73 Å
MCS + rigid confidence	2.27 Å	45 %	−0.396	1.73 Å
MCS + flexible confidence	1.41 Å	74 %	+0.065	1.20 Å
MCS + flexible, 5 conformer starts	**0.86 Å**	**88 %**	**+0.264**	**0.71 Å**

Three things follow. The **torsions rescue the confidence criterion**: with a conformation to fix, the rigid search of the 2019 paper is worse than not refining at all, which is the failure the protocol was built to expose; the same objective with the rotatable bonds free is the best method available. The **conformer generator is what makes it work** — five generated starting conformations are worth 0.55 Å on the median pose, better on 15 of 16 sets, and the combination beats plain substructure alignment on 16 of 16. And the **displacement guard has little left to prevent**: with these starts no set is worse than 1.52 Å, a 1 Å guard costs 0.6 Å, and 2 Å is indistinguishable

from none.

What none of it does is replace knowing the conformation: +0.264 against the crystal poses' +0.334. The conformational residue is 0.71 Å, and that is where the remaining error lives.

Using it

In the application, *Tools → Align by prediction confidence…* gained a **turn the rotatable bonds as well** tick box, with the number of starting conformations and the strain allowance beside it. From the command line:

```
cmf predict model.cmf new.sdf --align-variance --align-flexible --align-conformers 5
```

The force field acts as a prior on the conformation — it may rise ten kcal/mol above where it started before being penalised, one-sided so that the aligner does not quietly become a conformer optimiser — and the displacement limit now applies to the **rigid motion only**, since a turning bond is not a translation and charging it as one would suppress the very thing this adds.

Charges and per-atom properties

The electrostatic field needs charges: Gasteiger (the default), MMFF94, or **read from the file**, which is what you want with the Sutherland sets and their MNDO ESP-fit charges, or with anything you computed yourself.

Beyond that, any per-atom number can become a field of its own. `prop:<name>` reads the atomic weights from a property of the molecule — a quantum charge, a Fukui index, a local softness, an atomic polarisability, even a per-atom embedding — and that field then behaves exactly like a built-in one, with its own attenuation, its own mixing coefficient, its own co-field, atom contributions and Shapley values.

Two conventions are read, in this order:

1. a property set on **every atom**, which RDKit writes to SDF as `atom.dprop.<name>` and reads back automatically;
2. a **molecule property holding one number per atom**, separated by spaces, commas or newlines, which is how most quantum-chemistry packages dump such a table.

The hydrogens are the trap. A file without 3D coordinates is embedded on loading and hydrogens are added, so a list written for the heavy atoms alone no longer matches the atom count. The program says so rather than silently mis-assigning: `--prop-missing error` (the default) stops, `--prop-missing zero` gives the unlisted atoms zero weight. The next section shows the whole route from an ORCA run to fields and descriptors, from inside the program or from the command line.

Quantum-chemical properties with ORCA

Tools → Quantum-chemical properties with ORCA… runs ORCA over the loaded molecules and brings two kinds of number back, to the two places the model keeps them.

Per atom, and so available as `prop:` fields: Mulliken, Löwdin, Hirshfeld and CHELPG charges, the HOMO and LUMO population of every atom, and — when the N+1 and N−1 states are computed as well — the condensed Fukui indices f^+, f^-, f^0, the dual descriptor $f^+ - f^-$, the local softnesses and the local electrophilicity. **Per molecule**, and so available as descriptors: the total energy, the HOMO and LUMO energies and their gap, the dipole moment, and from the ionised states the ionisation potential, the electron affinity, the chemical hardness, the softness, the chemical

potential and the electrophilicity index.

That split is the whole point of the feature. A Fukui index is a statement about one atom, so it belongs in a field, where a co-field can then say *where* a nucleophilic centre helps; the gap is a statement about the molecule, so it belongs in the descriptor block, where it cannot pretend to have a position.

The dialog asks for

- **the ORCA program** — found by itself at the standard location, in `$ORCA_EXE` or on the PATH, and otherwise given here;
- **the level of theory** — the ORCA keyword line, which may be edited: `r2SCAN-3c` (fast, with its own basis) is the one to start with, B3LYP-D3/def2-SVP the usual hybrid, and anything else ORCA accepts as a single point will do;
- **which molecules** — all of them; only those that do not carry the properties yet, which is the default, so a set extended with new compounds costs only the new ones; the selected ones; or the training set;
- **the Fukui indices and the global reactivity numbers**, which need the N+1 and N−1 states and so cost three calculations per molecule instead of one;
- **optimising the geometry first**, off by default — see the warning below;
- **how many molecules at a time**, and, where the ORCA build has MPI, how many cores per job. The Windows build is the serial one: it has no MPI helpers, `%pal` fails outright, and the way to use a many-core machine is to run several molecules at once, which is what the first box is for. The dialog disables *cores per job* when it sees a serial build;
- **memory per core**, a **time limit** per job, after which that molecule is given up on rather than holding up the set, and the **working directory**.

Before you press *Run* the dialog says what the run will cost: the number of molecules times the number of ORCA jobs each, and the wall time that comes to at the chosen parallelism. After the first run that estimate is no longer a guess — the time every job took is recorded in the working directory, and the dialog quotes the measured mean. For scale: a molecule the size of methanol at `r2SCAN-3c` takes a few seconds per job, a drug-sized one at B3LYP-D3/def2-SVP a couple of minutes, so a hundred of the latter with their Fukui indices is a few hours with four running at a time.

The dialog also states the charge and multiplicity it read off the structures — the formal charge of the molecule, and the lowest spin state its electron count allows — so an unintended anion or an odd electron is visible before the hours are spent, not afterwards in a column of nonsense.

Three things make a long run survivable. Every job is **cached** under a key over the molecule, its charge, the level of theory and the geometry, so *Cancel* — which kills the ORCA processes themselves, not merely the bookkeeping — loses only the jobs in flight, and starting again continues from where it stopped. A molecule that **fails** is named in the log and skipped, never fatal to the set. And the working directory keeps every input and output, so any number can be traced back to the calculation that produced it.

When the run finishes, the new quantities appear as tick boxes — the per-atom ones in *Per-atom property fields*, the per-molecule ones in *Descriptors (per molecule)* — and the program offers to write the annotated set as an SDF. Write it: the file carries the properties with it, as atom property lists and molecule properties, so opening it later brings them all back without running ORCA again. *Tools → Save the data set with its properties (SDF)…* does the same at any time.

> **About the geometry.** *Optimise the geometry first is off by default and should normally stay off. ORCA optimises each molecule on its own, which moves every atom: a series that was superimposed is not superimposed afterwards, and a model trained on the old poses no longer refers to the molecules in memory. With the box off, the properties are computed exactly where the molecules stand, which is what a field model needs. If you do optimise, align the set again before building anything.*

The same computation runs from the command line as `examples/orca_atom_properties.py` (`--help` lists its options), which reads an SDF and writes the annotated one. The script and the dialog share the module `cmfqsar/orca.py`, so the numbers are the same either way.

Before you model: `cmf diagnose`

Every expensive mistake recorded in this manual was visible before the model was built, in a number that takes seconds to compute. The corrosion benzenes were modelled with five fields and an automatic frame, and atom-type counts — no geometry at all — beat them. The ACE set is predicted as well by a handful of two-dimensional descriptors as by the whole 3D model. A thrombin model in the exposure study scored +0.059 with its attenuation pinned to the top of the grid, which is what a mis-specified model looks like from outside. None of these is subtle; they are simply never asked, because asking means running four baselines by hand.

```
cmf diagnose data.sdf --charges file
```

The command answers five questions and prints the numbers behind each.

Is there anything to model? The size of the set and the spread of the activity. Twenty molecules across half a log unit cannot support a conclusion, whatever q^2 comes out, and the command says so and stops.

Can a frame exist at all? The number of chemotypes, the size of the common Bemis–Murcko framework, and whether the molecules are actually superimposed. That last test is asked of the shared framework atoms rather than of the centroids, because `EmbedMolecule` leaves every molecule centred on the origin: a set of independently generated conformations has a centroid spread of 0.02 Å and looks perfectly aligned while every molecule points a different way. The test is also allowed to decline. A framework that is a small symmetric ring matches a dozen ways and two rings near the origin may coincide by luck, while a framework that is a small part of the molecule may match a *different* copy in each one — measured that way an author-aligned Sutherland set read as 9 Å apart. So a coincidence within half an ångström is a verdict on its own, a large deviation is a verdict only when the framework is a fair share of the molecule, and otherwise the answer is "this cannot tell — look at it".

Does the frame carry signal? The aligned model against the two routes that need no frame: the atom-type count kernel, which has thrown the geometry away entirely, and the rotation-invariant field autocorrelation, which keeps it without a pose. If neither is beaten, the alignment is decoration and the cheaper route is also the honest one.

Is the geometry decoration? The same comparison against descriptors alone — what the file carries plus the handful the program computes.

Was the model specified sensibly? Whether the chosen attenuation sits on the edge of its grid, which means the search wanted a value it was not offered.

A fifth route, the topomer rule, is scored where the molecules can be cut at their framework. It covers only those molecules, so the aligned model is refitted on the same ones for the comparison: a subset against the whole set is not a comparison, and it flatters whichever route happens to drop the hard compounds.

Two things the command protects you from

The activity as a descriptor. Most files store the activity as a molecule property, so the descriptors available usually include the quantity being predicted. Handed it, the descriptor baseline returns $q^2 = 1.000$ and the advice becomes "use descriptors" for every set in the world.

Columns named as the activity are dropped, and so is anything correlating with it above 0.95 — but the leak is rarely a copy. The corrosion files carry the potency as `activity`, again as `mwila_pooled_raw`, and again as two laboratory-adjusted variants; excluding the first two still left a descriptor baseline at $q^2 = 0.99$, because the remainder reconstruct it between them. No per-column test can see that, so the outcome is judged as well: a descriptor baseline reaching $q^2 \geq 0.95$ on a set of this size is not chemistry but the measurement itself, and the command says so, names the nearest columns with their correlations, and refuses to recommend descriptors.

Searching harder and calling it a better model. The five routes do not carry the same number of hyper-parameters — the field routes search six attenuations and five ridges, the others search ridges only — so scoring each at the best point of its own grid tilts the comparison toward whichever searched hardest. Every route therefore chooses its hyper-parameters inside each training fold and is scored outside it. On ACE this costs the aligned model 0.05 and the topomer route 0.07 while leaving counts, descriptors and the invariant kernel where they were, which is the size and the shape of the bias it removes.

Reading the answer

Each comparison is reported with a paired bootstrap interval over the molecules, because on twenty to thirty molecules a difference below about 0.1 in q^2 is not a difference. The last line recommends a route, and it should be read as a prompt rather than a verdict: it says which road to take and which question the set has not answered, never that a model is good.

One wording deserves care. When the command recommends `invariant` it often means *the frame is not worth the trouble* rather than *the invariant kernel predicts better* — on DHFR the aligned model is nominally ahead, +0.659 against +0.641, and the recommendation follows the interval rather than the point estimate. And `aligned, with the frame checked` means the frame carries signal and the molecules really are superimposed, so a frame exists whoever made it; several Murcko scaffolds only mean the program could not have built it, not that it is wrong.

Does the advice work?

A tool that reports five baselines is useful whatever its last line says, but advice that has never been followed is only a rule its author liked. `benchmarks/run_diagnose.py --validate` measures it the only way that counts: hold out 30 % of a set, run the diagnosis on the remaining 70 % alone, build the model it recommends, and score that model on the held-out molecules against every other route — thirteen sets, three splits each.

The advice named the best of the five routes on **21 of 39 splits**, and averaged over all of them the route it chose fell **0.096** in q^2 short of the best available. Always taking the aligned model — the default this command exists to question — names the best route on 12 splits and costs 0.201. Aggregated per set, which is the honest unit because three splits of one set are not three measurements, following the advice costs **0.105 less in q^2 than always aligning** (95 % bootstrap interval −0.259 to +0.019), better on six sets, worse on four, indistinguishable on three.

The shape of that distribution matters more than its mean. The advice is a **hedge, not an edge**: where it loses it loses a little — 0.02 to 0.17 on AChE, BZR, THR and the thermolysin set, all of them well-aligned series where the aligned model was always going to be right — and where it wins it wins a lot, 0.44 to 0.70 on the three corrosion sets, where the automatic frame is meaningless and the aligned model collapses. The interval crosses zero, so on thirteen sets this is not a proven improvement in the average case. What it is, is insurance against the failure the chapter opened with, bought at a price of a few hundredths.

Building a model

Choosing the fields

The five physico-chemical fields — electrostatic, steric, hydrophobic, H-bond acceptor, H-bond donor — are the default and reproduce the published CMF models. The steric field is special: its width is scaled by the van der Waals radius of each atom (α/R^2_{vdW}, Eq. 10–11 of the 2013 paper) rather than being shared.

Continuous indicator fields (CIF, 2015 paper) add one field per atom type, with weight 1 on atoms of that type and 0 elsewhere. Tick *Continuous indicator fields* and set the minimum number of training molecules an atom type must appear in (3 by default) to be given a field of its own. Three typing schemes are available: `sybyl` (the types of the original scripts, read from MOL2 or assigned), `element`, and `element+aromatic`. CIFs are the most descriptive field set the program has, and on several of the benchmark sets they match or beat the physico-chemical fields on their own.

`prop:<name>` fields are added by naming them in the per-atom property row of the model panel, or in `--fields`.

GRID-like probe fields

Goodford's GRID describes a molecule by the energy a small probe — a water, a methyl, a carbonyl oxygen — would feel at every point around it. That is a field in exactly this program's sense, and it says something the built-in fields do not: not "there is hydrophobic character here" but "a probe placed here would be bound". *Probe fields* in the model panel (`--probes water,carbonyl`) adds one such field per probe. Five are defined:

name	what it stands for	R (Å)	well (kcal/mol)	charge	hydrogen bonding
`water`	water, OH_2	1.70	0.15	0	both, weakly
`methyl`	a methyl group	2.00	0.10	0	neither
`carbonyl`	a carbonyl oxygen	1.60	0.20	−0.40	accepts
`amine`	an amine nitrogen	1.75	0.16	+0.40	donates
`cation`	a generic +1 cation	1.50	0.10	+1.00	neither

Those four numbers are the whole definition of a probe: the radius and well depth enter the dispersion term through the usual combining rules, the charge the electrostatic term, and the hydrogen-bond character decides which atoms it is complementary to — a donor probe feels acceptor atoms and not other donors. The parameters sit in the usual GRID and Tripos ranges rather than being taken from a published GRID table, which matters if you ever want to compare numbers with GRID itself rather than shapes. Adding a probe is one entry in `PROBES` in probes.py; the fitting, the kernel, the drawing and both interfaces read that dictionary, so a new one appears everywhere at once.

The obstacle is arithmetic, not conceptual: a Lennard-Jones term is not a Gaussian, has no closed-form overlap, and diverges at the atom, so the integral of its square does not exist. The way through is to fit each atom's contribution by a few Gaussians, which keeps every kernel in the program a closed form over atom pairs. Three things had to be measured before that worked, and each one shaped what is implemented:

- **The repulsive wall cannot be carried.** Fitting a curve that is +30 kcal/mol at 2.5 Å and −0.13 at 3.5 Å with a handful of Gaussians produces coefficients of $\pm 10^{7}$ that cancel, and a worst-case error larger than the well it is meant to represent. The wall is a truncation artefact in every grid implementation anyway, and this program already carries it — the steric field *is* a Gaussian van der Waals envelope — so a probe field carries the **attraction** only.
- **A well at a radius is a shell**, and a shell is not a sum of Gaussians centred on the atom. Each term is therefore capped at its own optimum, which makes it monotone in r, and those a sum of Gaussians represents naturally.
- **The r^{-6} tail is not Gaussian either**, and is negligible beyond 6 Å, so the energy is tapered to zero between 5 and 7 Å exactly as GRID's own cut-off does.

What comes out reproduces the attraction to about 15 % in the volume-weighted norm the kernel sees (23 % on a realistic worst case), and the kernel it implies agrees with numerical integration of the fields themselves to $4 \cdot 10^{-7}$ — a test asserts both. Each atom carries K Gaussians (four by default), so a probe field costs K^2 passes over the distance matrix where an ordinary field costs one; three is the cheap setting and five the careful one.

Measured against the fields it would join, same molecules, same α and λ grids, 10-fold CV:

set	n	the five fields	+ water probe	probes only
GPB	66	+0.539	+0.539	+0.406
THER	76	+0.538	+0.543	+0.428
THR	88	+0.672	**+0.699**	+0.665
ACE	114	+0.647	+0.648	+0.607

So: adding a water probe to the five fields is worth between nothing and 0.027 q^2, and probe fields on their own are worse than the fields they would replace. That is the honest result and it was the expected one — CoMSIA replaced GRID-like energies with Gaussian similarity indices for much the same reason, and this program's fields are closer to CoMSIA than to GRID by design. What a probe field is actually *for* is interpretation: its co-field is drawn, exported and decomposed like any other, and it says where a water or a carbonyl oxygen would be bound, which is a more directly actionable statement than "more acceptor character here".

Two limitations are worth stating plainly. The directional part of GRID's hydrogen-bond term — the angles at the donor and at the acceptor's lone pairs — is **not** represented; the radial part is, broadened to 0.6 Å because a sharper shoulder is not representable by the shared basis. And alignment by field similarity refuses to run on a probe field rather than silently ignoring it, because the gradient of the expansion is not implemented.

Weighting the fields by solvent exposure

A kernel that integrates over all space weights an atom buried inside a molecule as heavily as one on the outside, although a binding partner can only touch the second. That is the claim the surface methods are built on, and `cmfqsar/exposure.py` tests it the cheap way: every atomic weight

multiplied by its atom's relative solvent accessibility, which keeps the kernel in closed form and keeps the co-field.

It does not help, and on the larger sets it mildly hurts. Across sixteen FEP sets the median paired gain is −0.002 in q^2 and the weighting wins on eight of sixteen — a coin flip — while the one large win is a model whose attenuation was pinned to the grid edge. On the eight Sutherland sets every variant is worse, and the softened one that looked best on the FEP sets has an interval entirely below zero. The effect also shrinks as the sets grow: on DHFR (361) and COX2 (282) the change is −0.004 and +0.003, and the movers are the two smallest sets. An effect that vanishes with more data is noise.

The reason is geometric and settles the general question. Measured by group — a heavy atom and its hydrogens counted as one, because atom by atom a methyl carbon is 93 % occluded by its own hydrogens and ethanol reads as two-thirds buried — the median group exposure is 0.37 to 0.42 and **0 to 3 % of groups are buried** in either family. A ligand of thirty heavy atoms has no interior. Surface-based representations are solving a problem that molecules of this size do not have.

Descriptors that are not fields

Everything above describes a molecule as a function of space. Some things are not functions of space and never will be: a measured logP, a molecular weight, the medium or temperature a series was recorded in, a count of rotatable bonds, an energy from a quantum calculation. A model that cannot see them has to pretend the activity depends on geometry alone, which for an adsorption potency or a permeability is simply false.

The *Descriptors (per molecule)* row of the model panel (`--descriptors A,B`) offers every numeric molecule property the file carries — and hides itself when the file carries none besides the activity, which is why it is invisible on most of the example sets; the computed descriptors of the next section are always there. Ticking some adds one more block to the kernel:

$$K(A,B) = \sum_f h_f c_f \hat{K}_f(A,B) + h_d c_d \, \mathbf{z}(A) \cdot \mathbf{z}(B)$$

where **z** holds the descriptors of a molecule, standardised with the **training set's** mean and standard deviation so that a quantity in kilojoules cannot outvote one in log units, and h_d is a mixing coefficient like a field's (the *h* box beside the row, `--descriptor-h`; zero switches the block off without removing it).

Three properties of the block are deliberate:

- **It is scaled, not cosine-normalised.** Dividing by the two norms, as the fields are divided, would send a molecule sitting at the training mean in every descriptor — an unremarkable molecule — to a zero row and a zero diagonal, out of the model altogether. Instead the whole block is multiplied by a constant that puts its mean diagonal at 1, so ten descriptors do not outweigh the geometry merely by being ten.
- **The statistics belong to the training set** and are stored in the saved model. Standardising a test molecule by its own set would score two test sets on two different scales, which is the mistake the field scaling avoids as well.
- **The model stays linear in its inputs**, so the decomposition rules survive. A prediction splits into a field part and a descriptor part, and the descriptor part splits further into one number per descriptor, which is what the program reports.

A molecule missing a descriptor is refused by default (`prop_missing`), because a model that silently substitutes the mean is telling you nothing about that molecule while looking as confident as ever; with the `mean` policy it is substituted and the molecule simply contributes nothing through that descriptor. A descriptor that is constant over the training set is reported as inert rather than dividing by a spread of zero.

The obvious trap, stated plainly. A descriptor derived from the activity — the activity itself, a transform of it, a value fitted to it — will dominate the model and make the geometry look worthless, while q^2 goes up. The program cannot detect this, so it reports the mean absolute contribution of every descriptor beside the model statistics:

```
Non-field descriptors (mean |contribution| per training molecule):
  clogp                          0.0779
  mw                             0.0176
  tpsa                           0.0123
```

If one descriptor carries most of the prediction, the model is a descriptor model wearing a 3D coat, and the honest response is to report it as one.

```
python cmf_cli.py build series.sdf --activity pIC50 --descriptors clogp,mw,tpsa
```

Descriptors the program computes itself

A file that carries no numeric tags has nothing to offer the row above, so four descriptors are always available, computed by RDKit from the structure: `calc:mw` (molecular weight), `calc:clogp` (Wildman–Crippen logP), `calc:tpsa` (topological polar surface area) and `calc:rotatable` (rotatable bonds). They have a row of their own in the model panel, *Computed descriptors*, and **all four are off by default**.

Those four are shortcuts, not the whole list. **Any** of RDKit's descriptors — some two hundred of them — can be asked for by its own RDKit name, typed into the box at the end of that row (with completion) or passed on the command line: `calc:NumHAcceptors`, `calc:FractionCSP3`, `calc:BalabanJ`, `calc:LabuteASA`, `calc:qed`. Names are matched case-insensitively, and a name that does not exist is refused before the fit with the nearest ones that do — `calc:NumHDonor` answers *did you mean NumHDonors?*.

The default is deliberate, and so is the warning that belongs with it. These are two-dimensional quantities: they can lift a model that has a size or lipophilicity trend in it, and they can also carry a model that then owes nothing at all to its geometry. Measured on four Sutherland sets, same α and λ grids, 10-fold CV:

set	n	five fields	+ the four computed	the four alone
THR	88	**+0.671**	+0.634	+0.620
GPB	66	**+0.534**	+0.525	+0.494
ACE	114	+0.657	+0.707	**+0.712**
THER	76	**+0.509**	+0.496	+0.480

Read ACE carefully: four numbers that know nothing of the geometry predict better than the whole 3D model. That is not an argument for using them — it is a warning about that data set, and it agrees with what the alignment-free baseline already said about ACE elsewhere in this manual. On

the other three sets they cost a little.

So use them to *test* a model rather than to improve it: if adding them changes nothing, the geometry is carrying the model; if they take it over, the 3D part was decoration. The contribution table printed beside the statistics is what tells you which happened.

Attenuation, regularisation and the grids

Two hyper-parameters matter. α sets how far each field reaches: large α makes the fields sharp, so molecules must match closely, while small α makes them diffuse. λ (or *C*, ε, ν for the SVM learners) sets the regularisation. Both are chosen by cross-validated grid search, and the whole grid is shown in the *Hyper-parameter search* tab — look at it, because an optimum at the edge of the grid means the grid was too small.

```
python cmf_cli.py build data.sdf --alphas 0.1,0.3,0.5,1,2,5 --lams 1e-3,1e-2,0.1,1,10
```

Refinements, all optional: continuous Brent minimisation over λ as in the original R scripts (`--lam-search brent`), a Nelder–Mead search over α (`--alpha-search nelder-mead`), coordinate descent or centred kernel alignment for the mixing coefficients `h_f` (`--optimize-h`, `--mkl`), and a grid over the bandwidth of a nonlinear kernel layer (`--gammas`).

Two conventions decide how the field kernels are combined, and they change the numbers: *Normalise kernels* uses cosine-normalised field kernels, and *Scale fields* rescales each field kernel to unit mean diagonal before mixing. Untick the second to get the raw-kernel convention of the papers (`--raw-scale`).

The learners

Task	Methods	Key
Regression	kernel ridge regression with bias; Gaussian process (the same fit plus predictive variances); ε-SVR; ν-SVR; the ranking SVM, which fits the order and not the values	`krr`, `gp`, `svr`, `nusvr`, `ranksvm`
Binary classification	C-SVC; ν-SVC; Gaussian-process classifier by the Laplace approximation	`svc`, `nusvc`, `gpc`
One-class / screening	1-SVM; distance to the centre of the actives; distance to the k-th nearest active; one-class Gaussian process; distance from the principal subspace (kernel PCA)	`ocsvm`, `ocmean`, `ocknn`, `ocgp`, `ockpca`

Kernel ridge regression is the method of the papers and the default. The Gaussian-process view adds a posterior variance per prediction at no extra cost, which is what the confidence-based alignment of *Preparing the data* minimises. The Laplace GP classifier adds calibrated probabilities, a posterior variance and model selection by marginal likelihood — it scores no better than C-SVC on the benchmark sets (mean difference 0.001 in Matthews correlation), so it is not the default, but its probabilities are honest and its evidence picks hyper-parameters without holding anything out.

Fitting the order instead of the values: the ranking SVM

A regression model is asked for the activity and judged on the squared error, which takes the numbers at their word. Often they do not deserve it. Potencies pooled from several papers carry an offset per source that no descriptor explains, a series measured in two assays carries two scales, and an IC50 read off the end of a dose–response curve carries a bound rather than a value. In all three the *order* survives what the values do not, and a model fitted to the values spends its capacity on the part that is noise.

`--method ranksvm` fits the order. For every pair of molecules whose activities differ it asks that the more active one score higher by a margin, and pays for every pair where that fails:

$$\min \tfrac{1}{2}\|w\|^2 + C\sum_p \xi_p \;\text{ with }\; w\cdot(X_i - X_j) \geq 1 - \xi_p$$

over the pairs $p = (i, j)$ with $y_i > y_j$ — an ordinary support-vector classification on the *difference* of two molecules' fields. The kernel of two such differences is a sum of four ordinary kernel entries,

$$K_{\text{pair}}((i,j),(k,l)) = K(i,k) - K(i,l) - K(j,k) + K(j,l)$$

so the whole thing is built from the kernel matrix the program already has, with no new fields and no new parameterisation.

The fitted model is still a dual expansion over molecules. The weight vector is $w = \Sigma_p\, a_p\, (X_{i_p} - X_{j_p})$, so collecting the terms gives one coefficient per *molecule* and the score of a new one is $K(x, \text{train})\,\mathbf{a}$, exactly as for every other learner here, with no bias. Everything that reads a linear CMF model — co-fields, atom contributions, Shapley values, the pharmacophore — keeps working. What changes is what they mean: the co-field of a ranking model says where more field *promotes* a molecule, not what it adds to the activity.

What is given up is the scale. A ranking score has no units and cannot be compared with an activity; only its order means anything. The program therefore reports Spearman's ρ, Kendall's τ and the share of pairs ordered correctly, and reports no r^2 or RMSE, which would be arithmetic on a quantity that does not exist. The hyper-parameter search selects on ρ for the same reason.

```
python cmf_cli.py build data.sdf --charges file --method ranksvm --Cs 0.1,1,10 \
    --rank-gap 0.5 --rank-groups paper
```

Three options shape the pairs. `--rank-gap` drops the pairs whose activities differ by less than it: a pair separated by less than the experimental error is a coin toss, and asking a model to call it is how a ranking model learns noise. Half a log unit suits pooled literature potencies and zero suits one careful assay. `--rank-groups PROP` forms pairs only inside a group — the same paper, assay or series — which is what makes the method worth having on pooled data, because an offset between two sources then cancels instead of being compared across. `--rank-pairs` caps the number of pairs: the solver works on the kernel *between pairs*, so the cost is quadratic in that number, and the default of 4,000 already asks for 128 MB. The order of a set is heavily over-determined and a few thousand comparisons fix it, so the cap is a bound on cost rather than a loss of information.

When it is worth it, measured

`benchmarks/run_ranking.py` asks the only question that matters — does fitting the order beat fitting the values — on the eight Sutherland sets with five random 70/30 splits each, every model choosing its own hyper-parameters on its own criterion inside the training part.

On clean data it does not. Where the activities are one assay's numbers and mean what they say, Spearman's ρ on held-out molecules is **+0.713 for kernel ridge and +0.689 for the ranking SVM**, ridge ahead on five of the eight sets. That is the expected answer and worth stating plainly: a regression model is already a ranking model, since the order of its predictions is all a ranking model ever produces, and it had the magnitudes to learn from as well. Fitting the order discards how far apart two molecules are, which on honest data is information.

On pooled data it does, but only with the groups. The second experiment gives the same molecules to three synthetic sources — assigned *by scaffold*, because an offset that travels with the chemistry is what a model will learn as if it were an effect of the chemistry, while an offset assigned at random is simply averaged away — shifts each source by 1.5 standard deviations, and asks for the order of the true activity, which no model ever saw:

	KRR	ranking	ranking by source
same-source pairs ordered correctly	72.5 %	71.6 %	**74.3 %**
Spearman ρ over all held-out pairs	+0.574	+0.565	**+0.663**

A ranking model that is not told about the sources does no better than ridge: it is handed the same corrupted comparisons and dutifully learns them. Told which molecules share a source, it beats ridge by **0.089 in ρ**, on six of the eight sets, because the cross-source pairs — the only ones the offsets corrupt — are never formed. Note that it wins on the *global* order too, although it was never shown a cross-source comparison: it does not need one, since it fits a single scoring function over the fields and all that grouping did was stop it fitting an offset.

So the honest summary is narrow and useful: **the ranking SVM is for data whose numbers are not comparable across their sources, and its value is** `--rank-groups`**.** On a single clean assay, use kernel ridge.

Nonlinear kernel layers and multiple kernel learning

A layer on top of the CMF kernel makes the model more expressive than an overlap integral: a Gaussian layer, $\exp(-\gamma\, d^2)$ with $d^2 = K(i,i) + K(j,j) - 2K(i,j)$, or one or two arc-cosine layers, the kernel of an infinitely wide neural network (Cho & Saul 2009). The price is interpretation: the model is then no longer linear in the field, so co-fields and atom contributions are unavailable. The similarity decomposition stays valid, and the **Shapley tab** decomposes the prediction instead, which is why it exists.

Multiple kernel learning chooses the mixing coefficients h_f rather than leaving them at 1, either by coordinate descent on the cross-validated error or by centred kernel alignment (`--mkl align|alignf`), which is a closed-form, training-only criterion.

Validation

- **k-fold cross-validation** (10 by default) gives q² and RMSEcv. Folds are shuffled unless you ask for the interleaved folds of the original scripts (`--cv-no-shuffle`).
- **A test set** gives R²pred and RMSEpred. Both variance conventions are reported, because they differ for sets whose test activities span a narrow range — the steroid benchmark is the classic example, where the ten test affinities cover 1.9 log units and R²pred computed against the test variance is negative while the predictions are good.

- **External (nested) cross-validation** (`--nested`, *External CV*) wraps the whole model selection in an outer loop and reports q²ex and RMSEcvex. This is the number to quote: the ordinary q² has seen the hyper-parameter choice.
- **y-randomisation**, a **bootstrap** interval and the **applicability domain** are the three sections below.
- **Leave-one-chemotype-out** folds, for a set of several scaffolds, never leave a molecule's own scaffold-mates in the training part. Random folds flatter such sets badly — by 0.11 to 0.47 q² on the five chemotype benchmarks here — and the program reports both, together with an alignment-free fragment baseline that the aligned model has to beat.

Is it chance? y-randomisation

A cross-validated q² says how well the model predicts; it does not say whether a model of the same shape could have reached that number from activities that mean nothing. In a congeneric series, where kernel values above 0.97 are normal, that is a real possibility, and it is the first thing a referee asks about a 3D-QSAR model.

The **shuffles** box of the model panel's *Validation* row (`--y-scramble N`) shuffles the activities N times, refits and cross-validates each time, and reports where the real q² falls among the scrambled ones — the mean and spread of the shuffled values, the largest of them, how many standard deviations above them the real one lies, and the permutation p-value.

```
python cmf_cli.py build data.sdf --alphas 0.5 --lams 0.1 --y-scramble 100
```

Two properties of the test are worth knowing, and the program states both rather than letting the number mislead:

- **A permutation test cannot report a p below 1/(N+1).** Ten shuffles bottom out at p = 0.091, which is above every conventional threshold *however good the model is*. The program says so — "no shuffle came close, but 10 of them cannot support a p below 0.091" — instead of reporting a chance-correlation warning that the data do not support. Use 19 shuffles for $p < 0.05$ and 99 for $p < 0.01$.
- **The hyper-parameters are held fixed** by default, so the test asks whether the *fit* can be reached by chance. Tick **redo the search** beside the shuffles box (`--y-scramble-select`) for the stricter question — whether the whole procedure, model selection included, can be. It matters: on the steroid set with a 3 x 3 grid, the best scrambled q² rises from −0.08 to **+0.32** once the search is allowed to pick the luckiest grid point on noise, while the real value stays at +0.78. The kernels do not depend on the activities, so they are computed once and the extra cost is one cross-validation per grid point per shuffle.

Where does the model apply?

A kernel method always returns a number. For a molecule unlike anything in the training set, that number is the training mean in disguise, and nothing in q² or R²pred reveals it.

The **applicability domain** tick of that row (`--domain`) measures how close each molecule outside the training set is to it, **in the kernel the model itself uses**: the cosine similarity of their fields, averaged over the *k* nearest training molecules (3 by default). The threshold is not a number chosen out of the air — it is taken from the training set's own spread, the 5th percentile of the same quantity computed among the training molecules, so a series of close analogues gets a tight domain and a diverse set a loose one. Molecules below it are listed by name in the log, their

similarity appears in the data-set table's **AD** column — which sorts like any other, so the extrapolations can be brought to the top — and the exported predictions carry both the value and a yes/no verdict, so that an extrapolation stays marked after it leaves the program.

On the steroid benchmark, nine of the ten test compounds fall inside and one — 4-pregnene-3,11,20- trione — falls outside. On BZR, training on the largest chemotype and scoring the others, every molecule of the training chemotype is inside and 29 of 40 molecules from other chemotypes are outside, which is the behaviour one wants from a domain estimate.

The field-space distance is reported alongside the similarity, because the two can disagree: a molecule that is a scaled-up version of a training molecule points the same way (high cosine) but sits far away (large distance).

Is the whole library inside? The two-sample test

The applicability domain answers the question one molecule at a time. A screening run asks a different one — is this library drawn from the same population as the training set, or is the whole exercise an extrapolation? — and averaging a per-molecule verdict does not answer it: two sets can have identical nearest-neighbour similarities and still occupy different regions.

Every set of molecules has a **mean field**, the mean of its feature maps, which for CMF is itself a function of space. The maximum mean discrepancy (Gretton *et al.*, JMLR **13** (2012) 723) is the distance between two of them:

$$\mathrm{MMD}^2(X,Y) = \overline{K(x,x')} - 2\,\overline{K(x,y)} + \overline{K(y,y')}$$

— three averages over blocks of the kernel matrix and nothing else. The unbiased estimator leaves the self-similarities out of the two within-set terms, which is why it can come out slightly negative: that is how "no difference at all" looks once the bias has been removed, not a defect. The null distribution comes from permuting the labels, which costs three block means per shuffle because the kernel is computed once.

```
python cmf_cli.py compare library.sdf --model model.cmf --permutations 999
python cmf_cli.py compare seriesA.sdf seriesB.sdf --charges file
```

With `--model` the library is compared with that model's training set, in that model's own fields, reusing the cross-kernel the prediction computes anyway. The p-value is never zero: it is one plus the number of shuffles that matched or beat the observed value, over one plus the number of shuffles, because a finite number of them cannot support a claim of impossibility.

The witness says where the difference lies, which is the actionable form of a p-value. It is the function whose expectation separates the two sets — a molecule's mean kernel value against the first set minus its mean against the second — evaluated at every molecule: most negative for the compounds that make the library unlike the training set, most positive for those most typical of it. The command prints both ends and `--out` writes the whole column. On the Sutherland benzodiazepine set against the COX2 one, the witness names nordazepam and delorazepam as the most typical benzodiazepines and the COX2 sulphonamides as the least — without being told that there are two targets.

Is the test any good?

A test is worth nothing unless it is calibrated as well as powerful, and `benchmarks/run_structure.py` measures both halves. Twenty random half-splits of each of the eight Sutherland sets, where there

is nothing to find: **5.6 % are called significant at the 5 % level**, which is what a correctly calibrated test does and the permutation construction guarantees. All twenty-eight pairs of *different* targets, where there is everything to find: **100 % are called significant, median z = +171**. The test is calibrated where there is nothing and saturated where there is.

How certain is that number?

A difference of 0.02 in q^2 is quoted all over the benchmarks in this repository. Some of those differences are real and some are the particular molecules they were computed over.

The **bootstrap** box of that row (`--bootstrap N`) resamples the (observed, predicted) pairs N times and reports a percentile confidence interval for q^2 and for R^2_{pred}. It is the uncertainty of the *statistic* given these predictions, which is exactly what decides whether two settings differ. On the 21 training steroids, $q^2 = 0.768$ carries a 95 % interval of [0.47, 0.89] — wide enough to settle any argument about a hundredth.

`--bootstrap-refit N` adds Cramer's bootstrap as well: resample the training set with replacement, refit, and score the molecules that were left out of the bag. That one measures the variability of the fit rather than of the statistic, and it is the more pessimistic of the two. The kernel is never recomputed — a resample is an index into the matrix — so a hundred refits cost about as much as one cross-validation.

How certain is one prediction? Conformal intervals

The bootstrap says how firm a figure of merit is. It says nothing about the molecule in front of you, and neither does anything else the program has offered so far: the Gaussian-process posterior spread is the spread of a model that believes its own kernel, which for molecular fields it should not.

Conformal prediction supplies the missing guarantee and asks almost nothing in return. Take the absolute errors the model makes on molecules it did not fit, let q be their $(1-\alpha)$ quantile — the $\lceil (n+1)(1-\alpha) \rceil$-th smallest, the +1 being what makes the result finite-sample rather than asymptotic — and report $\hat{y} \pm q$. Then, as long as the calibration molecules and the new one are exchangeable, the truth lies inside at least $(1-\alpha)$ of the time, **whatever the learner, whatever the kernel, however wrong the model is**. A badly specified model does not lose its coverage; it pays in width, which is the honest currency.

```
python cmf_cli.py build data.sdf --charges file --conformal 0.9 --save model.cmf
python cmf_cli.py predict model.cmf new.sdf --out predictions.csv
```

`--conformal` calibrates on the errors of the cross-validation the program is running anyway, so it costs nothing and adds no step, prints the half width, and — when there is a test set — the coverage it actually achieved on molecules that were never part of the calibration. The rule is stored in the `.cmf` file, so `predict` prints an interval beside every number and writes `lower` and `upper` columns into the CSV. `--conformal LEVEL` on `predict` re-reads the stored errors at a different level; the errors are kept, not just the width.

What it costs, measured

Calibrating on out-of-fold errors and then predicting with the model refitted on everything is not exactly exchangeable — the errors come from models fitted on nine tenths of the training set and the predictions from a model fitted on all of it — so the guarantee becomes an approximation, and the size of the gap is an empirical question. `benchmarks/run_conformal.py` answers it on the eight Sutherland sets, ten random 70/30 splits each, with the width in units of the activity (whose

standard deviation averages 1.36):

nominal	out-of-fold (`cv`)	held-out third (`split`)	GP posterior
80 %	**80.3 %**, width 2.28	80.1 %, width 2.59	77.7 %, width 3.00
90 %	**89.9 %**, width 3.10	90.1 %, width 3.61	83.1 %, width 3.85
95 %	**95.1 %**, width 4.09	94.9 %, width 4.04	86.9 %, width 4.59

The cheap variant holds its level to a few tenths of a per cent at every level, and is **narrower** than the exact split version, which pays for its held-out third by fitting on less data. The Gaussian-process interval, which is what one would otherwise reach for, misses its level at all three — and misses it unpredictably: across the eight sets its 90 % coverage ranges from 63 % to 99 %, while every conformal rule stays between 89 and 96 %.

Normalised conformal — dividing each error by the distance of its molecule from the training set, so that the interval widens where the model extrapolates — is available as `--conformal-normalise`, and on these sets it buys nothing: it over-covers (93.4 % at a nominal 90 %) and is a third wider. Softening it until it stops over-covering turns it back into the plain rule. The error on these series simply does not grow with distance from the training set strongly enough for the normalisation to pay, which is itself worth knowing about one's data.

Two limits worth stating. The guarantee is **marginal**, not conditional: 90 % of molecules are covered, not 90 % of *every kind* of molecule, so a subclass the model handles badly can be covered less often than that while the overall figure holds. And a level needs molecules to support it — 19 calibration errors for 95 %, 9 for 90 % — below which the program says the level is not reachable rather than returning its widest error and calling it an interval.

Regression only: conformal classification produces prediction *sets* rather than intervals, and is not implemented.

What the model thinks the groups are: clustering the kernel

Every grouped validation in this manual splits the molecules by their Bemis–Murcko scaffold, and the reasoning is sound: random folds put a molecule's own scaffold-mates in the training part and flatter a mixed series badly. But it rests on an assumption nobody checks — that a scaffold is what makes two molecules alike *to the model*. The model does not see scaffolds. It sees fields, and its notion of similarity is the kernel.

Spectral clustering partitions that similarity directly. The kernel is already a similarity matrix, so there is nothing to build: form the normalised Laplacian, take the eigenvectors belonging to its smallest eigenvalues, put the rows on the unit sphere and run k-means there (Ng, Jordan & Weiss 2002). Groups that are not linearly separable in field space are separable in those coordinates, which is the point of the construction. How many groups is decided by the **eigengap** unless it is given, and the gap is reported beside the answer so that a weak one can be seen for what it is.

```
python cmf_cli.py cluster data.sdf --charges file
python cmf_cli.py cluster data.sdf --clusters 5 --out clusters.csv
```

The command prints the clusters, the spectrum they came from, the activity of each, and the **adjusted Rand index** between them and the scaffold groups — the number that says whether the usual grouping groups what the model actually keys on.

The two partitions do not agree, and it matters

Measured on the eight Sutherland sets, with both partitions cut into the same number of groups — holding out one of two clusters removes half the training set and holding out one of nine does not, so a difference in granularity would be read as a difference in difficulty:

	q^2
random 10-fold	+0.508
leave one **scaffold group** out	+0.271
leave one **field cluster** out	+0.057

and the adjusted Rand index between the two partitions averages **+0.140**, which is weak agreement: what the chemist calls a series and what the kernel finds similar are largely different things. On BZR they are different enough that holding out a field cluster takes q^2 to −0.740.

The consequence is uncomfortable and worth stating. This manual recommends the leave-one-chemotype-out figure as the honest one, and against random folds it is — 0.271 against 0.508. But it is not the hardest honest test available: holding out a group of the model's own similarity costs another 0.214, because a scaffold group leaves molecules the kernel still considers close in the training part. Quote the scaffold figure, by all means; just do not read it as the worst case.

Where the validation appears

The three checks of *Building a model* — y-randomisation, the bootstrap intervals and the applicability domain — are set in one place and reported in another, which is worth stating plainly because there is no button and no separate window for them.

Setting them: the *Validation* row at the foot of the model panel's second group, directly under *External CV*. Three controls: **shuffles** (0 = off), **bootstrap** (0 = off), and the **applicability domain** tick. They are settings, not actions — choose them before pressing **Build model**, and they run as part of that build, after the fit and the test prediction.

Reading them: the *Statistics* tab of the results panel, where they are appended below the usual figures under a heading *Validation*. The names of the molecules that fell outside the applicability domain go to the *Log* tab instead, with their similarity values, so that the statistics stay a summary and the list of offenders sits with the rest of the commentary.

Three cases produce no validation block at all, and each says why rather than leaving a blank:

- all three controls at their defaults — nothing was asked for;
- a classification or one-class model — y-randomisation and the bootstrap are implemented for regression only, and the block says so;
- the autocorrelation kernel source — the applicability domain is measured in the field kernel, which that source does not provide.

Statistics reported

Regression: r^2 and RMSE for the fit, q^2 and RMSEcv for the cross-validation, R^2pred and RMSEpred for the test set in both variance conventions, q^2ex and RMSEcvex for the external CV. Classification: accuracy, balanced accuracy, sensitivity, specificity, Matthews correlation, AUC and the ROC curve, each for the fit, the cross-validation and the test set. Screening: AUC, enrichment factors at 0.5, 1, 5

and 10 %, the best achievable enrichment at each of those fractions, BEDROC and RIE.

Saving and reusing

File → Save model... writes a `.cmf` file (joblib) holding the field set, the kernels' definition, the fitted coefficients, the training molecules, every setting needed to predict later, and the record of how the training set was prepared that *What has to match between fitting and predicting* describes. `cmf_cli.py predict model.cmf new.sdf --out predictions.csv` applies it, after checking the new molecules against that record; the GUI loads it with *File → Load model...* and writes the record's summary to the log.

Reading the model

A CMF model is readable in a way a grid model is not, because the prediction can be taken apart exactly. The 2019 paper states three decomposition rules, and the program implements all three.

Co-fields

The co-field is the model as a function of space:

$$C_f(\mathbf{r}) = \sum_j a_j \, s_j \, X_{fj}(\mathbf{r})$$

a weighted sum of the training molecules' fields with the fitted dual coefficients as weights, each divided by the norm $s_j = K_f(j,j)^{-1/2}$ of its own field when the kernels are normalised, which they are by default. The prediction of any molecule is then the overlap of its own field with the co-field, field type by field type. Press **Compute co-field** and the program evaluates it on a grid around the training set; choose the field with the *Field* selector and set the *iso level* slider, which is a percentage of the largest absolute value rather than an absolute number, so the surfaces stay meaningful when you change field.

What the lobes mean: a positive co-field lobe is a region where *more* of that field raises the predicted activity. For the electrostatic field that means positive charge is favoured there, for the steric field bulk, for the hydrophobic field lipophilicity, and for an indicator field it means an atom of that type is welcome at that position.

Two further maps use the same grid, for the molecule currently shown:

- the **molecular field** $X_f(\mathbf{r})$ of that molecule, in the same colours. The two share a colour scale, so only the drawing style tells them apart: the style box is `auto` by default, which draws this one as a mesh while a co-field is on screen and solid when it is the only surface shown;
- the **overlap** $C_f(\mathbf{r})\, X_f(\mathbf{r})$, whose integral *is* that field's contribution to the prediction (Eq. 7), so the map shows where the prediction comes from.

File → Export current co-field (Gaussian cube) writes a cube file readable by PyMOL, VMD or Jmol, which is the way to put a co-field into a publication figure.

Contributions

The *Contributions* tab answers two questions about the selected molecule:

- **Which training molecules** make its prediction, as $a_j\, K(X, X_j)$ — the similarity decomposition, rule 2. A prediction resting on one or two training molecules is a warning.
- **Which of its atoms**, as the atom contributions of Eq. 14 of the 2015 paper, drawn in the 3D view as red and blue spheres whose size is the magnitude. This is the most directly actionable map the program produces: it says *this* substituent is helping and *that* one is not.

Shapley values

The *Shapley* tab decomposes a single prediction into additive contributions by the Shapley construction, over **atoms**, over **substituents of the common scaffold**, or over **field types**, with second-order interaction indices available. It serves two purposes. With the linear kernel it

reproduces the atom contributions exactly, which is a useful check. With a nonlinear kernel layer, where co-fields do not exist, it is the only decomposition left — and it still works, because it needs nothing but the ability to evaluate the model on subsets.

A pharmacophore out of the co-field

Tools → Extract pharmacophore from the co-field... turns the co-field into a discrete, typed pharmacophore: the favourable lobes of each field become feature points of the matching kind (positive and negative centres, hydrophobe, donor, acceptor), each with a tolerance radius taken from the size of its lobe, and the unfavourable steric lobes — or the pocket, when a protein is loaded — become exclusion volumes. The result is a table and a PDB-style file of feature spheres that other software can read, and it draws in the 3D view; untick *pharmacophore* in the viewer controls to clear it.

This is the bridge from a CMF model to a screening query: the model is continuous and needs an alignment, the pharmacophore is discrete and does not.

What a candidate is missing

Tools → Field difference from a reference molecule... subtracts two molecules' fields on the co-field grid and reports, per field type, the cosine between them, the distance, and how much of the reference's field the candidate **fails to reproduce** against how much it **adds**. The difference map draws as its own isosurface, and the table ranks the whole data set by field similarity to the reference. In a lead-optimisation reading: the "missing" column is what to put back, the "added" column is what may be causing trouble.

Two endpoints at once: the selectivity co-field

A separation problem, or any selectivity, is not one property but two: an extractant has an affinity for americium and an affinity for europium, and what matters is the difference. A model of the ratio works, but it throws information away. Each affinity has its own measurements, often on partly different compounds; a ligand that binds both metals strongly looks, to a model of the ratio, exactly like one that binds neither; and the ratio cannot say *which* of the two affinities a substituent moved.

Tools → Model two endpoints (selectivity co-field)... fits both endpoints on the kernel of the model you have already built, with one bias per endpoint and a task covariance that lets them borrow strength from each other:

$$y_t(X) = \sum_j c_{jt}\, K(X, X_j) + b_t,$$

where the coefficients of the two endpoints are tied through a single correlation ρ. At ρ = 0 the endpoints are modelled separately and share only the kernel; as ρ rises they borrow from each other; at ρ = 1 the two sets of coefficients coincide and the model could only predict a *constant* selectivity, so 1 is excluded (and ρ may be negative, for endpoints that move in opposite directions). The dialog searches ρ and the ridge by cross-validation, and the cross-validation holds out whole **molecules**, never single measurements: with a molecule's first endpoint in the training set and its second in the test set, the correlation would hand the answer over.

Because each endpoint's prediction is linear in the field with its own coefficients, each endpoint has a co-field, and so does their difference — the **selectivity co-field**

$$C_A(\mathbf{r}) - C_B(\mathbf{r}) = \sum_j (c_{jA} - c_{jB})\, X_j(\mathbf{r}),$$

a map of where more of a field raises the first endpoint more than the second. Only a model linear in the field can produce it; it is available here for the same reason co-fields are available at all, and it is refused for a model with a nonlinear kernel layer.

Once the fit is made, an **Endpoint** box appears in the *Maps* row of the control strip with three entries — the first endpoint, the second, and their difference — and every map, export and contribution table then refers to the chosen one; a cube exported while the selectivity is chosen carries it in its file name. The two endpoints are two numeric properties of the data set, so the second is imported with *File → Import activities from CSV...* or read from a column of the file. Molecules that carry only one of the two are kept: they inform the endpoint they have, and through the correlation the other. Each endpoint needs at least five molecules.

What the chosen ρ means is itself a result. On the corrosion-inhibitor series of the benchmark (`benchmarks/results/multitask/table.md`) the cross-validation chose ρ = 0 on the two largest series and a non-zero ρ of either sign on the small ones, where it gained at most 0.04 in q^2: an endpoint that already has its own molecules borrows little from another, and the ρ the dialog reports in the log is the measurement that says how much.

Chemical space maps

Tools → Chemical space map... builds a two-dimensional map of the set from the CMF kernels. Fields are not vectors, so Generative Topographic Mapping cannot be applied to them directly; the program performs kernel PCA on the CMF kernel — which is functional PCA, and reproduces feature-space distances exactly — and runs the ordinary EM fit of GTM in those coordinates. The out-of-sample projection stays a single kernel row, so new molecules can be placed on an existing map.

Two projections, and they answer different questions. The *Projection* box chooses between the GTM landscape and the **kernel PCA score plot**, which draws the functional principal components themselves — by default the first two, with any other pair available. The GTM spreads the set over a square grid, which is what makes a landscape possible, but it is a nonlinear mapping: a tight cluster is pushed apart to fill the grid and no distance on the map can be read as a distance between molecules. The score plot is the linear projection the GTM is fitted on, so its distances are the kernel's own as far as the two eigenvalues reach, and each axis is labelled with the share of the kernel variance its component explains — on the steroid set, 32.9 % and 15.1 %. Ask the map *where the activity lives*; ask the score plot *whether this series is one cluster or three*, and how much of the kernel a flat picture of it can honestly show. The score plot needs no EM fit, so it also costs a fraction of the time.

From the responsibilities follow the **activity landscape** (responsibility-weighted mean activity per node), the **class landscape** (the fraction of actives) and the **data density**, which is a picture of the applicability domain; nodes with no support are left blank. The landscape is a model in its own right: on the thrombin training set it reproduces the activities with R^2 0.92.

Two cautions. Use cosine-normalised kernels, or the leading components follow molecular size. And a CMF kernel between molecules that share no frame is meaningless, so a map spanning several series would partly be showing alignment quality — unless you use the pose-free count kernel of *Working without an alignment*, which separates three targets' ligands cleanly with no alignment at all.

Classification and virtual screening

Setting up a classification task

Choose *Binary classification* as the task. The classes come either from a threshold on the activity (with *active below* for data where small is potent, such as an IC50 in molar units) or from a class property in the file. The learners are C-SVC, ν-SVC and the Laplace Gaussian-process classifier; the *ROC* and *Scores* tabs replace the scatter plot.

```
python cmf_cli.py build examples/sutherland/AChE.sdf --charges file \
       --task classification --threshold 6.84 --method gpc --Cs 1,10,100
```

One-class models: screening with actives alone

Often there are no reliable inactives, only a set of known actives, and the question is whether a new molecule belongs to the same class. That is one-class classification, applied to CMF by Karpov et al. (2011). Five descriptions of the active class are implemented:

Method	What it measures
`ocsvm`	the 1-SVM: a hyperplane in feature space separating the actives from the origin
`ocmean`	distance to the mean embedding of the actives, the centroid of the fields
`ocknn`	distance to the k-th nearest active
`ocgp`	the one-class Gaussian process, optionally scored by mean over standard deviation
`ockpca`	reconstruction error against the principal subspace of the actives (kernel PCA)

On the sets measured here the 1-SVM is the **most fragile** of the five: the centroid, k-NN and kernel-PCA descriptions are as good or better and have one parameter fewer to go wrong. Try more than one; they disagree in interesting places.

Decoys and the figures that matter

A screening model is only as meaningful as the background it is tested against. *File → Add decoys (background set)…* or `--decoys` loads presumed-inactive molecules used **only** for evaluation, never for fitting. `examples/make_decoy_sets.py` builds property-matched decoy sets, which is the honest kind: decoys matched to the actives in molecular weight, logP, rotatable bonds and formal charge, so that the model cannot win by noticing size.

Two traps the program handles explicitly:

- **A decoy must not be rejected for its pose.** `--align-decoys` places every decoy in the frame of the training set by maximising its field similarity, with several random restarts, and `--decoy-conformers N` generates N conformations of each decoy and keeps the one that poses best (the file's own conformation is always among the candidates).

- **AUC saturates.** With cross-target decoys almost any method reaches 0.95 and above, which tells you nothing. Read the **enrichment factors** at 0.5, 1, 5 and 10 % instead, which the program reports beside the maximum achievable value at each fraction, and **BEDROC/RIE** (Truchon & Bayly), which weight the top of the list.

A measured caution from the benchmarks in this repository: field similarity on these tasks performs about as well as ECFP4 fingerprint similarity, and cross-target decoys are trivially separable by either. If a screening result looks spectacular, check what the decoys were.

Conformations: the fourth dimension

A 3D-QSAR model represents each molecule by one conformation, and the answer depends on which one. 4D-QSAR (Hopfinger et al. 1997) replaces the choice by an ensemble and averages over it. *Model → Conformations* switches between the two; **single is the default and is untouched**, and the tests assert that by comparing its fields against those built before ensembles existed.

In CMF the averaging needs no new mathematics. A field is a sum of atom-centred Gaussians and the kernel is bilinear in the fields, so the mean field of an ensemble is itself a field — its atoms are the atoms of every conformation with their weights multiplied by that conformation's weight — and the kernel of two ensembles comes out as

$$K(A,B) = \sum_k \sum_l w_k^A w_l^B \, K\left(X_A^{(k)}, X_B^{(l)}\right)$$

the mean-map kernel of the two conformational distributions, positive definite by construction. A test checks that identity against the pairwise kernels of the conformations. The model stays linear in that mean field, so **co-fields, the decomposition rules, Shapley values and the complementarity scores all keep working**; contributions arrive per atom per conformation and fold back onto the atoms by summation.

Tools → Generate conformations... builds the ensembles with ETKDG and MMFF94s, keeping the pose a molecule already has as its first conformation, so an earlier alignment survives. A multi-conformer file can be read instead, with `--ensemble-by name` grouping the records of one compound. Weights are uniform or Boltzmann from the stored energies, with kT and a cap on the number kept.

Where the other conformations go is a choice of its own, and not a small one: what matters is not where a conformation sits relative to its own molecule but where it sits relative to the rest of the series.

`--conformer-align`	**Where each generated conformation is put**
`self`	on the molecule's own pose, over all heavy atoms — the default, so nothing changes silently
`core`	on its own pose, over the common substructure with a reference only, so the scaffold stays where the published alignment put it and only the flexible part moves
`reference`	on a reference molecule of the series, through that same common substructure, which puts every conformation of the series in one frame
`field`	as `reference`, then refined by maximising the normalised CMF kernel — the 2013 paper's criterion, applied per conformation

Multi-instance learning (`cmf_cli.py mil`) goes further: instead of averaging, the model chooses which conformation of each molecule to use, alternating between fitting and selecting. Use it when you believe one conformation is the bioactive one and the others are noise; use the ensemble when you believe the molecule is genuinely flexible.

Working without an alignment

This chapter is the one to read if your molecules came from SMILES, if your series has several scaffolds, or if you simply do not trust the superposition you have. The program offers three frame-free routes, and the scripts in `benchmarks/` measure all of them.

Why the frame is the problem

An alignment is an assumption, and a wrong one is expensive. Measured on the nine targets that occur both in the curated Sutherland collection and in the automatically aligned Mittal collection — same compounds, same activities, same fields and grids — the curated superposition gives a mean q^2 of **+0.597** and the automatic one **+0.295**. Worse, the gain over a baseline with no geometry at all (atom-type counts) is +0.25 on the curated files and **−0.06** on the automatic ones: on five of the nine targets, fields in an automatic frame predict *worse* than counts. An alignment that is not trustworthy is not merely imperfect, it is actively harmful.

The pose-free count kernel

As α → 0 an indicator field tends to the count of atoms of that type, so the kernel tends to the dot product of the atom-type count vectors and nothing of the geometry survives. The program implements that limit in closed form (`chemotypes.count_kernel`), which is exact rather than approached, costs one pass over the atoms, and is 10 to 34 times faster than the Gaussian route. It is the **baseline** any aligned model should be made to beat, and it is what the leave-one-chemotype-out workflow reports beside the aligned model.

One subtlety, because it is easy to get wrong: the normalisation must be applied to the joint kernel, not field by field. A field normalised on its own tends to 1 whenever both molecules contain the type at all, so the counts cancel and the baseline loses the very descriptors it is built from.

The rotation-invariant autocorrelation kernel

Geometry without a frame. For each pair of field types the program computes the smoothed field autocorrelation

$$\Phi_{fg}(d) = \sum_{l}\sum_{m} w_{lf} w_{mg} \exp\left(-\frac{(|\mathbf{r}_l - \mathbf{r}_m| - d)^2}{2\sigma^2}\right)$$

as a function of interatomic distance d, which is invariant under rotation and translation by construction.

Where the two molecules are. That equation describes *one* molecule: both sums run over its own atoms, which is what makes it an autocorrelation. It is a descriptor, not a kernel — one number per channel c, a channel being one pair of field types at one distance — and the kernel is the cosine-normalised dot product of two such vectors:

$$K(A,B) = \frac{\sum_c s_c^2 \, \Phi_c(A)\, \Phi_c(B)}{\|s\,\Phi(A)\| \; \|s\,\Phi(B)\|}$$

with s_c the scale channel c inherits from its block: the reciprocal root mean square of that block over the training molecules, times the block's mixing weight. This is the opposite construction to

the overlap kernel of chapter 2, which is defined on two molecules from the start and has no finite feature vector at all — its feature map is the field itself. That reversal is also what decides how each model can be read afterwards: a dual expansion over a feature map that is a function of space is itself a function of space, which is the co-field, while a dual expansion over channels is a vector of channel weights, which is the co-pattern.

The $d = 0$ channel reproduces the CMF self-overlap exactly when $\sigma^2 = 2/\alpha$, so the descriptor contains the ordinary kernel's diagonal information and adds the distance profile. Select *Autocorrelation* as the kernel source in the model panel (`--kernel autocorrelation`, with `--ac-sigma` and `--ac-dmax`), and the model fits, cross-validates, saves and predicts exactly as any other — only the co-field, which needs a frame, is unavailable.

On the chemotype benchmarks the invariant kernel recovers most of the value of the geometry without a frame (0.172 against 0.189 for the aligned model and 0.072 for counts), and on automatically aligned files it frequently beats the aligned model outright — on six of the nine shared targets, including thrombin at +0.452 against +0.021.

Two points, and what that misses. The descriptor counts pairs of atoms at a distance, and a mirror image has every pair at exactly the same distance, so this kernel cannot tell a molecule from its enantiomer. It is blind for the same reason to the orientation of one field with respect to another — whether the donor sits clockwise or anticlockwise from the bulk, going round the ring. The quantity that is not blind is the signed volume of three displacement vectors, and `--ac-chirality` adds it as a second block of channels, the signed triple products of three field types resolved over `--ac-chirality-shells` radial shells (four by default). The block is appended rather than mixed in, so the two-point part is unchanged to the last bit, and it is scaled separately because a volume carries the cube of a length.

Measured, it does not pay for itself. `benchmarks/run_chirality.py` asks the question in the regime the alignment-free kernel exists for — leave-one-chemotype-out on the five chemotype-split sets, with three regularisations by three distance grids chosen inside the training part of every fold, and the number of shells kept out of the search so the signed channels never get more freedom than the kernel they extend:

	two-point	signed, 1 shell	signed, 2 shells	signed, 4 shells
mean q^2 over the five sets	+0.063	+0.082	+0.095	+0.108
paired difference	—	+0.019	+0.033	+0.045
95 % interval over the sets	—	[−0.051, +0.113]	[−0.031, +0.121]	[−0.028, +0.121]
better on	—	2 of 5	3 of 5	3 of 5

The mean moves the right way and the interval contains zero at every setting. More to the point, **one set supplies the whole of it**: at four shells COX2 gains +0.187 and the other four sets average +0.009 between them, with AChE and BZR made worse. On five sets that is not a result, and the channels cost 16 % more columns.

The experiment those sets cannot do. None of the five contains an enantiomeric pair, so the kernel's blindness costs it nothing there and the table above is about something else. The same script therefore builds a set where handedness *is* the activity: twenty molecules, each with its exact mirror image — one conformation reflected, not a second embedding, so a pair differs in handedness and in nothing else and their two-point channels agree to one part in 10^{15} — with the

activity `(1 − w)·composition + w·handedness`, the handedness read as the CIP descriptor rather than the SMILES tag, which is not the same label from one scaffold to the next.

At `w = 0` the model works: q² +0.97. At `w = 1`, where the answer is the handedness and nothing else, **no variant learns it** — two-point −0.391, and the signed channels −0.34 to −0.37 however many shells they are given and even when scaled up to carry half the kernel.

And the reason is measurable. A kernel method carries what it learnt from one molecule to another along a direction in feature space. The signed channels of a molecule and its own mirror image are exact negatives — that is what the block is for, and a test asserts it — but between *different* molecules they share nothing: the cosine of the signed block is **−0.016 for two molecules of the same handedness and +0.016 for two of opposite handedness**, a difference of 0.03 against a spread of 0.55. Even within one scaffold family, amino acids differing only in the side chain, the difference is 0.10 against a spread of 0.48. Handedness is not a direction in this space, so there is nothing for a linear model to transfer, and the two null results above are one null result seen twice.

That is the verdict: the block is correct and does what it says about a molecule and its mirror image; it does not put the handedness of two different molecules on a common axis, which is what a QSAR model needs. The channels stay available and off by default. Making them useful would mean a different construction — one anchored on the stereocentre rather than summed over the whole molecule — and that is a new item, not a parameter of this one.

Mixing coefficients over the channel blocks. The channels come in blocks, one per pair of field types, each scaled to unit root mean square before all of them enter one cosine-normalised dot product. The overlap kernel weights its fields with h_f; the three invariant sources (autocorrelation, surface, periodic) weight their blocks with `--block-weights` (library: `ac_block_weights` on `CMFModel`), one coefficient per block applied to the scaled features before the normalisation, so that a block's share of the kernel goes as the square of its weight and a weight of zero removes the block exactly. Three ways to set them: `balanced` gives every *kind* of block — the pairs within a group of fields, the pairs across two groups — the same share of the kernel however many blocks each kind has; `KIND=W,...` names the kinds' weights (`indicator=0.3,indicator*physchem=0.5`; unnamed kinds weigh 1); `align` or `alignf` derive one weight per block by centred kernel alignment of the per-block kernels (`align_block_weights`, the `align_mixing` of the fields carried over, with `--block-rho` the shrinkage of the joint solution), re-derived from the training part of every cross-validation fold, so the reported q^2 never sees the left-out activities. The groups default to the continuous indicator fields (`indicator`) against the rest (`physchem`); `--block-groups 'GROUP=field,field;GROUP=cif:*'` names others. The fitted weights are printed by kind and stored with the model (`block_weight_report`).

Molecules on a surface: the in-plane invariant kernel

A molecule adsorbed on a surface has one direction that is not like the others. Its height above the metal and its tilt are physics — they are what the relaxation computed, and they are most of what decides the adsorption energy. Where it sits in the plane and how it is turned about the normal are not: the slab is periodic and its origin is wherever the cell was cut. The overlap kernel cannot tell the two apart, so it demands a frame for a degree of freedom that has no preferred value, and the invariant kernel of the last section throws the height away along with the rest.

This kernel keeps exactly what is physical. It is the two-point distribution of the field weights over the **in-plane** distance ρ, resolved by the heights of both atoms:

$$\Phi_{fg}(\rho_k, z_i, z_j) = \sum_l \sum_m w_{lf} w_{mg}\, G_\sigma(\rho_{lm} - \rho_k)\, G_{\sigma_z}(z_l - z_i)\, G_{\sigma_z}(z_m - z_j)$$

$$G_s(u) = \exp\left(-\frac{u^2}{2s^2}\right), \qquad \rho_{lm} = |\mathbf{r}_l^{xy} - \mathbf{r}_m^{xy}|$$

It is therefore invariant to translation in the plane and to rotation about the normal, and to nothing else: tilt, height and the vertical order of the atoms all survive. Averaged over the direction of ρ it is, by the Wiener–Khinchin theorem, the in-plane power spectrum of the fields resolved by height — the molecule described in reciprocal space across the surface and in real space along the normal. The surface is at $z = 0$ with the normal along $+z$, which is how the adsorption scripts of `examples/corrosion_fe110` write their poses; nothing checks that convention, so a set of poses in some other frame will be described faithfully and uselessly.

The channels run over in-plane distances to 12 Å in steps of 0.75 Å (`--ac-dmax`, `--ac-sigma`) and heights from 1 to 6 Å in steps of 1 Å (`--surface-z ZMIN ZMAX ZSTEP`), the spacing of a channel being also its width on both axes. That is 17 × 6 × 6 = 612 channels for each pair of field types, 9,180 features on the five physico-chemical fields, which is why the blocks are scaled and may be weighted exactly as in the preceding section. Pairs of different field types are summed both ways round, as the three-dimensional descriptor counts them, and the $l = m$ terms are kept. Those channels are again a descriptor of one molecule, and the kernel of two of them is the dot product of the preceding section, over this section's channels:

$$K(A,B) = \frac{\sum_c s_c^2\, \Phi_c(A)\, \Phi_c(B)}{\|s\,\Phi(A)\|\ \|s\,\Phi(B)\|}, \qquad c = (f, g, \rho_k, z_i, z_j)$$

The command is `--kernel surface`; the model panel does not offer it.

What it is worth

On the benzimidazole corrosion node — 72 molecules relaxed on Fe(110) with the UMA machine-learned potential, the target their adsorption energy, five physico-chemical fields, KRR with α and λ chosen in a 10-fold inner cross-validation, 5 outer folds, 10 repeats:

model	q^2	RMSE (kJ/mol)
surface kernel, poses as the relaxation wrote them	**0.655**	31.2
surface kernel, the same poses shuffled in the plane	**0.655**	31.2
overlap fields, poses as written	0.505	37.4
overlap fields, poses registered in the plane by their common substructure	0.517	36.9
overlap fields, poses shuffled in the plane	0.480	38.3
rotation-invariant autocorrelation, which discards the height	0.109	50.2

Three things follow. The **invariance is demonstrated rather than asserted**: turning every pose by a random angle about the normal and sliding it up to 6 Å in the plane changes the model by

nothing at all, the second row reproducing the first to three decimals. **The height is what earns the score** — the fully invariant kernel differs from this one only in discarding *z*, and it falls from 0.655 to 0.109. And **the frame the overlap kernel wants is not worth building**: registering the poses laterally on the common substructure moves it from 0.505 to 0.517, a difference of nothing, while the kernel that was never given a frame stands 0.14 above the better of them.

There is no co-field, for the reason every invariant source has none: the kernel is not an integral of a product of fields, so the prediction cannot be turned back into a function of space. What it offers instead is the co-pattern — the weight of every channel of the fitted linear model, which names the pair of field types, the in-plane distance and the two heights that carry the signal, which `top_channels` ranks against the labels of `surface_channel_labels`. The measurements above come from `surface_kernel_eads.py`, in `examples/corrosion_fe110`, which writes its table beside itself.

Crystals: the periodic autocorrelation kernel

The same descriptor for a periodic structure, summed over the neighbour list of one cell with its periodic images out to a cutoff $r_{\max}$, per atom of the cell, with the uniform background of each field pair removed and the cutoff tapered:

$$\Phi_{fg}(d) = \frac{1}{N}\left[P_{fg}(d) \; - \; \frac{S_f S_g}{V}\, B(d) \right]$$

The first term is the tapered pair histogram of the two fields, over every atom l of the cell and every neighbour m inside the cutoff, periodic images included:

$$P_{fg}(d) = \sum_{l \,\in\, \text{cell}} \; \sum_{m \,:\, r_{lm} < r_{\max}} w_{lf} w_{mg}\, \omega(r_{lm}) \exp\!\left(-\frac{(r_{lm} - d)^2}{2\sigma^2} \right)$$

The second is what that histogram would be if the same total weight of each field were spread uniformly through the cell, which is the part of P_{fg} that reports only how dense the crystal is:

$$B(d) = \int_0^{r_{\max}} 4\pi r^2\, \omega(r) \exp\!\left(-\frac{(r - d)^2}{2\sigma^2} \right) \mathrm{d}r$$

Here N is the number of atoms in the cell, r_{lm} the distance from atom l to image m, $S_f = \Sigma_l\, w_{lf}$ the total weight of field f over the cell, V its volume, σ the width of the distance smearing (`--sigma`) and ω a half-cosine taper over the last `--taper` ångström. Both sums run over ordered pairs, so a pair of two different field types is counted both ways round, as the molecular descriptor counts it, and its background with it; the $l = m$ terms are kept, which makes the $d = 0$ channel the composition of the cell. As in the two preceding sections Φ describes one structure, and the kernel of two of them is the same cosine-normalised dot product over the channels:

$$K(A,B) = \frac{\sum_c s_c^2\, \Phi_c(A)\, \Phi_c(B)}{\| s\,\Phi(A) \|\; \| s\,\Phi(B) \|}, \qquad c = (f, g, d_k)$$

It is invariant to the choice of cell, origin and orientation, so a supercell is the same material. Averaged over orientations it is the sine transform of the product of the typed structure factors, which makes the pair products of the element indicator fields the partial structure factors a diffraction experiment measures.

The fields of a solid are built by `inorganic_fieldset`: electrostatic (oxidation states guessed from the composition, or the `oxi_states/charges` array of the file), steric (atomic volume),

electronegativity, one indicator field per element and one for the rest, and any per-atom array as `prop:<name>`. The command is `build-crystals` (an ASE-readable structure file, `--target KEY` or `--target-file CSV`, `--cutoff`, `--sigma`, `--taper`, `--no-background`, `--elements`, `--charges`, `--external-cv`, `--save`); in the library, `kernel_source="periodic"` on a `CMFModel` with a table from `crystal_table`. There is no co-field, but there is a co-pattern: the weight of every channel in the fitted linear model, which `top_channels` ranks by weight times spread, with the sign that says whether the pair at that distance raises or lowers the property.

On the Matbench `phonons` task (1,265 crystals, the top of the phonon density of states) the linear kernel on four fields -- electrostatic, steric, electronegativity and the inverse square root of the mass -- reaches a mean absolute error of 117 cm^{-1} on the Matbench folds, 106 with a Gaussian layer, against 324 for the dummy predictor and 29 for the best graph network; keeping the uniform background improves the 58-field model from 181 to 119 on this task, so the switch is worth trying on a property that depends on composition and density (the four-field model kept it too reaches 92). The element indicator fields dilute the joint kernel — seventeen hundred equally scaled element-pair blocks against the ten physico-chemical ones — and `--block-weights` is the remedy: with the background kept, the aligned weights take the 59-field model from 119 to 97; they do not make it better than the four physico-chemical fields alone. The full table and the co-patterns are in `benchmarks/results/matbench_phonons/table.md`.

Topomer fields

Cramer's topomer idea (topomer CoMFA, 2003), applied to fields. Cut the molecule at the bonds that attach its substituents to a common core and rebuild each substituent in a pose that is a **deterministic function of its topology**: canonical atom numbering, one embedding with a fixed seed, every rotatable torsion set to 180°, the attachment atom at the origin with its open valence along −x, and the remaining roll fixed by the principal axis of the heavy atoms. The same R-group then always lands in the same place, whatever molecule it came from — a test asserts the coordinates are identical, not merely close.

Two things make this fit CMF better than it fits CoMFA. There is no grid, so a fragment's field is just its atoms and a fragment–fragment kernel is an analytic number, which means the fragment library can be tabulated once and reused. And each attachment point keeps **its own region of space** (the substituent of position k is placed at $x = 100k$ Å, far enough that no Gaussian reaches the next site), so a molecule's kernel is exactly the sum of its per-position overlaps — which is what topomer CoMFA obtains by giving every position its own block of columns. A test checks that identity to 1e-9.

In the application it is *Tools → Replace molecules by their topomer substituents…*, and it is a **transformation of the data set** rather than a hidden switch: the molecules are replaced by their posed substituents, names, activities and the split are kept, and the data-set line says *[topomer substituents of …]*. That is the honest shape for it, because a topomer model's co-field and contributions live in the substituent frames and not around the original ligands. Reload the file to get the molecules back.

Choosing the core matters more than anything else, which is why the program now chooses it differently. The maximum common substructure of a chemotype mixture is not a scaffold — on GPB it comes out as `CCO`, fused into every ring system, and cuts nothing off at all, losing every molecule. The default is therefore the **Bemis–Murcko framework of the largest chemotype**, the scaffold of one series, with every molecule that framework comes cleanly off forming the series; on GPB that is `C1CCOCC1` and 60 of the 66 molecules. The common substructure is still available for a set you know is congeneric (*maximum common substructure of the series* in the dialog, `--core-mode mcs`), and you can always give the core yourself as SMILES or SMARTS. Check the cuttable fraction

the program reports either way: a core fused to the rest of the molecule cannot be removed without opening a ring, and what is left is not a substituent.

What the three routes are worth, measured

Every column below uses the same molecules, fields, α and λ grids and folds.

Collection	Series	Aligned q²	Topomer q²	The cut costs	The pose costs
Sutherland, curated frame	8	**+0.445**	+0.371	+0.058	+0.015
Mittal, automatic frame	31	+0.200	**+0.256**	−0.047	−0.009
Corrosion inhibitors, neutral	17	+0.248	+0.249	+0.061	−0.061
Corrosion inhibitors, 1 M HCl	17	+0.303	+0.292	+0.056	−0.045

The last two columns separate the two things the topomer construction does at once, by way of a control that keeps the same cut substituents at the same separated sites but leaves each one exactly where the alignment put it. The **pose rule is nearly free** — 0.015 q² against a careful superposition, and a *gain* of 0.045 to 0.061 where the frame is automatic, because a canonical pose is more consistent than an arbitrary embedded conformation. What costs something is the **cut**: separate sites discard the scaffold's own field and the geometry between one substituent and another, which is expensive when the arms' mutual placement is the pharmacophore (0.26 on thermolysin, 0.18 on COX-2) and free or better than free when the scaffold is constant and the decorations small (BZR gains 0.123, DHFR 0.071).

Practical guidance. With a trustworthy curated alignment, keep it: the geometry is worth about 0.25 q² over counts. With an automatic alignment or none, prefer the frame-free routes — topomer for a congeneric series with a clear scaffold whose R-groups come off at acyclic single bonds, and the invariant kernel for anything pooled, multi-scaffold or ring-fused. Always report the count baseline beside the model, so that the value of the geometry is visible rather than assumed.

Heterogeneous series, step by step

When a set holds several chemotypes with no one common substructure, the *Tools* → *Heterogeneous sets* submenu is a four-step workflow:

1. **Find the chemotypes** — Bemis–Murcko frameworks, optionally generic, with small groups merged so that leave-one-chemotype-out does not degenerate into leave-one-molecule-out.
2. **Align within each chemotype** — a superposition that is meaningful, because the members of one cluster do share a substructure.
3. **Place the chemotypes in one frame** — each cluster positioned by maximising field similarity to the others, which is the best that can be done when no common substructure exists.
4. **Leave one chemotype out, with the alignment-free baseline** — the honest validation, reported beside the count kernel. Random folds flatter such sets by 0.11 to 0.47 q² on the five chemotype benchmarks here, which is why this step exists.

Protein pockets

Getting a pocket on screen

File → Fetch protein from the PDB... downloads an entry from RCSB and caches it under `~/.cmfqsar/pdb`, with the dialog pre-filled with the identifier the data set names. *File → Load protein structure...* reads a local PDB or mmCIF file instead. The reader is deliberately independent of RDKit, because crystal structures carry alternate locations, missing elements and non-standard residues that a chemistry toolkit would refuse to sanitise, and none of that matters for drawing.

The free-energy benchmark sets in `examples/openff` and `examples/schindler` give their ligands in the frame of a reference crystal structure and name it per record, so **nothing has to be superimposed**: opening such a set brings its pocket with it from `examples/pockets/`, which holds the residues within 8 Å of any ligand, taken whole, without waters and without the co-crystallised ligand. The *protein site* controls above the 3D view then draw the residues within a chosen distance as wire frame or sticks, with hydrogens, waters and residue labels optional.

Complementarity as numbers

A pocket is a set of atoms, so the same parameterisation that builds a ligand's fields builds the protein's, and the overlap of two Gaussian fields has the same closed form on either side. The *Complementarity* tab (CLI: `cmf_cli.py complement`) turns that into four answers:

- **How complementary a molecule is to the pocket** — overlap integrals between ligand and pocket fields, exact and without a grid, each divided by the two norms so that it is a cosine in $[-1, 1]$ and does not grow with molecular size. Complementarity is not similarity: the electrostatic term carries a minus sign, because a positive ligand region should meet a negative protein region; the hydrophobic term carries a plus; the hydrogen-bond terms pair the donors of one side with the acceptors of the other. Steric overlap is reported separately as *burial*, because bulk against the protein is contact and clash at once, with the number of clashing heavy-atom pairs beside it.
- **Where on the molecule it holds** — the same quantities on a smooth van der Waals surface give a value per vertex, drawn as a green-to-red surface, with a correlation per field over the surface. Electrostatics uses potentials rather than Gaussian densities here, because the potential is what the other side feels.
- **Whether the co-field agrees with the pocket** — the model's own map against the protein, by sign agreement, cosine or correlation. This is the strongest external check a CMF model can be given: the co-field was fitted to activities and never saw the protein.
- **Which residues say so** — the same scores, resolved by receptor residue.

Which residues say so

An overlap integral is a double sum over atom pairs, so keeping the per-atom column sums instead of the grand total costs nothing and is exact: every term of the score above splits over the residues of the pocket and adds back to it. One cosine becomes an **interaction profile**, a table of residues against interaction types.

This is worth having because a global score cannot say which contact a pose has lost, and neither can the ligand RMSD — a pose can move very little and still give up the one interaction that

mattered. Across a series the profiles also average into a **consensus**: an interaction seen in most of the complexes is a property of the target, while one seen in a single complex is more likely a property of that ligand, which is the distinction a pharmacophore query needs.

On the ten CDK2 ligands of `examples/openff/cdk2.sdf` against their pocket, every molecule reaches 18 of the 53 residues, and the contacts conserved across the whole series are Glu81 donating to an acceptor, Leu83 accepting from a donor, and Ile10 and Leu134 by burial — the hinge and the hydrophobic floor, recovered without being told that either exists.

In the application, tick **by residue** beside *surface scores* before pressing *Compute*. The table under the Complementarity tab then lists every residue within 4.5 Å of the selected molecule, one column per interaction type, sortable by any of them, with the share of the series that shows each residue's interactions in the last column; the header says how much of the whole score those residues carry, the rest being the tail of the Gaussians over the residues behind them. Tick **contributions** beside *residue labels* in the protein-site controls and the 3D view writes each residue's total on its alpha carbon, green where it helps and red where it does not.

From the command line:

```
python cmf_cli.py complement data.sdf --structure pocket.pdb --residues --profile-out profiles.csv
```

`--residues` prints the strongest contacts of each molecule and the conserved pattern of the series, and scores every molecule against that consensus by cosine and by Tanimoto over the thresholded bits; `--reference NAME` compares against one molecule's profile instead of the consensus, `--residue-threshold` sets what counts as present, and `--profile-out` writes one row per molecule and residue.

Docking

Tools → Dock into the pocket... places the conformations of a molecule in the site by rigid-body optimisation. It is a convenience for getting a pose into the right frame, not a scoring function; the program's business is what happens after the poses exist.

Command-line reference

Everything the application does is available without it, which is what makes the program scriptable for a benchmark or a cluster run. The general form is

```
python cmf_cli.py <subcommand> <dataset> [options]
```

and `python cmf_cli.py <subcommand> --help` prints the full option list; what follows is the map, not every option.

Subcommand	What it does
`build`	build, cross-validate, test and save a model
`predict`	predict with a saved model, optionally aligning by confidence first
`hetero`	the heterogeneous-series workflow: cluster, align, place, leave-one-chemotype-out
`topomer`	model a series by its substituents, posed by rule
`map`	chemical-space map and activity landscape (kernel PCA + GTM)
`explain`	Shapley decomposition of one prediction, for any kernel
`mil`	multi-instance learning over conformations
`pharmacophore`	read a pharmacophore out of a fitted model's co-fields
`difference`	what one molecule's fields have that another's do not
`complement`	field complementarity with a protein pocket
`unalign`	give every molecule a random pose, the control experiment

The options of `build`, by group

Data. `--activity PROP`, `--activity-file TABLE`, `--test-file FILE`, `--test-fraction F`, `--ensemble-by KEY`.

Fields. `--fields electrostatic,steric,hydrophobic,hba,hbd[,prop:name]`, `--cif`, `--cif-scheme sybyl|element|element+aromatic`, `--cif-min N`, `--cif-types LIST`, `--charges gasteiger|mmff|file`, `--param rdkit|tripos|tripos-legacy`, `--prop-missing error|zero`.

Kernel. `--kernel fields|autocorrelation|surface`, `--ac-sigma S`, `--ac-dmax D`, `--surface-z ZMIN ZMAX ZSTEP`, `--ac-chirality`, `--ac-chirality-shells S`, `--block-weights W`, `--block-groups G`, `--block-rho R`, `--kernel-layer linear|gauss|arccos1|arccos2`, `--gammas LIST`, `--normalize`, `--raw-scale`, `--optimize-h`, `--mkl [align|alignf]`, `--mkl-rho R`.

Task and learner. `--task regression|classification|oneclass`, `--method krr|gp|svr|nusvr|svc|nusvc|gpc|ocsvm|ocmean|ocknn|ocgp|ockpca`, `--threshold T`, `--active-below`, `--class-property PROP`, `--oc-k K`, `--oc-components Q`, `--oc-lam L`, `--oc-score mean|mean_over_sd`, `--no-bias`.

Hyper-parameters. `--alphas LIST`, `--lams LIST`, `--Cs LIST`, `--epsilons LIST`, `--nus LIST`, `--lam-search grid|brent`, `--alpha-search grid|nelder-mead`, `--nm-start A,L`.

Conformations. `--conformations single|ensemble`, `--conformers N`, `--conformer-window A`, `--prune-rmsd A`, `--conformer-align self|core|reference|field`, `--conformer-align-ref I`, `--conformer-weighting uniform|boltzmann|property`, `--kt KCAL`, `--max-conformers N`.

Validation. `--folds K`, `--cv-no-shuffle`, `--nested`, `--nested-on train|all`, `--outer-folds K`, `--repeats R`, `--seed S`.

Screening. `--decoys FILE`, `--align-decoys`, `--decoy-restarts N`, `--decoy-conformers N`.

Alignment. `--align-train-variance [MAX_RMSD]`, `--align-train-rounds N`, `--align-train-folds K`.

Output. `--save model.cmf`, `--predictions out.csv`, `--json out.json`.

Worked examples

```
# indicator fields, a test split and nested cross-validation
python cmf_cli.py build examples/benzamides_synthetic.sdf --activity pActivity --cif \
       --alphas 0.3,0.5,0.8,1.2 --lams 1e-3,1e-2,0.1,1 --test-fraction 0.25 \
       --nested --save model.cmf

# predict, with the new molecules aligned by prediction confidence first
python cmf_cli.py predict model.cmf new.sdf --align-variance 2.0 \
       --aligned-out aligned.sdf --out predictions.csv

# the published protocol of the 2013 paper on a Sutherland set
python cmf_cli.py build examples/sutherland/THR.sdf --charges file --param tripos-legacy \
       --raw-scale --cv-no-shuffle --lam-search brent --nested --nested-on all

# a frame-free model: the rotation-invariant autocorrelation kernel
python cmf_cli.py build series.sdf --kernel autocorrelation --ac-sigma 0.75 --ac-dmax 12

# a topomer model with the core given explicitly
python cmf_cli.py topomer examples/benzamides_synthetic.sdf --activity pActivity \
       --core "c1ccccc1C(=O)N" --alphas 0.3,0.5,1 --lams 1e-2,0.1,1

# one-class screening against a property-matched background
python cmf_cli.py build actives.sdf --task oneclass --method ockpca \
       --decoys background.sdf --align-decoys --decoy-conformers 5

# a classification model with the Gaussian-process classifier
python cmf_cli.py build examples/sutherland/AChE.sdf --charges file --task classification \
       --threshold 6.84 --method gpc --Cs 1,10,100
```

Using it as a library

The library is the layer the application and the command-line tool are built on, and it is small enough to use directly in a notebook.

```
from cmfqsar import FieldSet, CMFModel, load_dataset
from cmfqsar.fields import AtomTable, build_fields, cif_types_in
from cmfqsar.model import grid_search

ds = load_dataset("aligned.sdf", activity_prop="pIC50")
fs = FieldSet.create(cif_types=cif_types_in(ds.mols, "sybyl", 5))
mf = [build_fields(m, fs, n) for m, n in zip(ds.mols, ds.names)]
train = AtomTable.from_molecules(mf)
model = CMFModel(fieldset=fs, method="krr")
grid_search(train, ds.activity, model, alphas=[0.3, 0.5, 1.0], lams=[1e-3, 1e-2, 0.1, 1])
model.fit(train, ds.activity, cv_folds=10, reuse_kernels=True)
print(model.stats.summary())
```

Decomposing one prediction, which works for the nonlinear kernels as well:

```
from cmfqsar.shapley import atom_shapley, field_shapley, scaffold_decomposition

res = atom_shapley(model, mf[0])                       # one value per atom
res = field_shapley(model, mf[0], interactions=True)   # one value per field type
res = scaffold_decomposition(model, mf[0], ds.mols[0], ds.mols[1:20])   # per substituent
print(res.prediction, res.baseline, res.table(), res.top_interactions(5))
```

The modules, one line each:

```
cmfqsar/fields.py       field parameterisation; FieldSet / MoleculeFields / AtomTable
cmfqsar/kernels.py      analytic field kernels, vectorised
cmfqsar/model.py        KRR / SVR / GP, cross-validation, grid search, nested CV, statistics
cmfqsar/cofields.py     co-fields and molecular fields on grids, isosurfaces, cube export
cmfqsar/dataio.py       reading SDF / MOL2 / SMILES, model persistence, exports
cmfqsar/align.py        3D embedding, MCS alignment, field-similarity refinement, unalign
cmfqsar/varalign.py     alignment by prediction confidence (GP posterior variance)
cmfqsar/invariant.py    rotation-invariant field autocorrelation kernels
cmfqsar/topomer.py      topomer fields: substituents posed by rule
cmfqsar/chemotypes.py   scaffold clusters, placement, grouped folds, the count kernel
cmfqsar/pharmacophore.py  typed features with tolerance spheres, from the co-field
cmfqsar/fielddiff.py    field differences against a reference molecule
cmfqsar/shapley.py      Shapley decomposition over atoms, substituents, field types
cmfqsar/gpc.py          Laplace Gaussian-process classification
cmfqsar/oneclass.py     centroid, k-NN, GP and kernel-PCA one-class descriptions
cmfqsar/conformers.py   conformational ensembles: generation, weighting, grouping
cmfqsar/mil.py          multi-instance learning over conformations
cmfqsar/decoys.py       posing a background set for screening evaluation
cmfqsar/gtm.py          chemical-space maps and activity landscapes
cmfqsar/protein.py      PDB / mmCIF reading, binding-site selection, frame check
cmfqsar/complementarity.py  pocket fields, complementarity scores, co-field agreement
cmfqsar/docking.py      rigid-body placement of conformations in the pocket
cmfqsar/gui/            the PySide6 application
```

The example data sets

Set	Contents
`examples/benzamides_synthetic.sdf`	109 aligned N-phenylbenzamides with a *synthetic* activity: the overlap of their fields with a hidden pseudo-receptor co-field plus noise (σ = 0.25). A model has to recover those zones; $q^2 \approx 0.87$ against a noise ceiling of 0.94. The set to learn the program on, because the right answer is known.
`examples/steroids_cbg.sdf`	the classic CoMFA benchmark: 31 steroids (Cramer, Patterson & Bunce 1988) with corticosteroid-binding-globulin affinities, in the corrected Gasteiger-group geometries, re-superimposed on the steroid skeleton of testosterone. Customary split 1–21 / 22–31, leave-one-out $q^2 \approx 0.79$.
`examples/sutherland/*.sdf`	the eight literature sets of the 2013 paper (Sutherland, O'Brien & Weaver 2004): ACE, AChE, BZR, COX-2, DHFR, GPB, THER, THR, with the aligned MMFF94s conformations as distributed, the published split, and the original Sybyl types and MNDO ESP charges as per-atom properties. Choose *Partial charges: file*.
`examples/mittal/*.sdf`	38 of the 40 benchmark sets of Mittal, McKinnon & Sorich (2009), built here from their 2D supplementary data by ETKDG and an MCS superposition, with each molecule's fit quality recorded in `align_rmsd`. The collection to use when you want to know how a method behaves on an *automatic* alignment.
`examples/openff/*.sdf`	21 congeneric series, 513 ligands, from the Open Force Field protein–ligand benchmark, each posed in its target's crystal structure; activity tag `pAffinity`.
`examples/pockets/*.pdb`	the binding sites of those series, trimmed and already in the ligands' frame, so they open together with the data.
`examples/thrombin_atom_properties.sdf`	a worked example of `prop:<name>` fields: per-atom quantum properties carried in an SDF.

`docs/benchmark.md` compares the program with Table 2 of the 2013 paper on the eight literature sets. `benchmarks/run_sutherland.py` reruns that comparison, `benchmarks/run_topomer.py` measures the alignment against the topomer rule (`--mittal` for the automatically aligned collection), `benchmarks/run_alignment.py` prices a curated alignment against an automatic one, and `benchmarks/run_corrosion.py` does the same for corrosion-inhibitor series.

When something goes wrong

The 3D view is black, or the program will not start. The viewer needs working OpenGL through PyOpenGL and pyqtgraph. On a remote desktop or in a virtual machine, software rendering is usually the fix (on Windows, a Mesa `opengl32.dll` beside the program; on Linux, `LIBGL_ALWAYS_SOFTWARE=1`). Everything except the 3D view works regardless, and the command-line tool needs no graphics at all.

q² is high but the test predictions are poor. Three usual causes, in order of frequency: the hyper-parameters were chosen on the same cross-validation you are quoting (run *External CV*); the set holds several chemotypes and random folds are flattering it (use leave-one-chemotype-out); or the alignment is doing work that will not transfer (compare against the count baseline).

q² is negative. The model predicts worse than the mean: check that the activity column is the right one and on a logarithmic scale, and that the molecules really share a frame. For R^2_{pred} on a narrow test set it may be an artefact of the variance convention, and the program reports both.

The optimum sits at the edge of the α or λ grid. Extend the grid. The *Hyper-parameter search* tab exists to make this visible.

A per-atom property is rejected. The hydrogen count almost certainly changed when the file was loaded. Write the property for all atoms including hydrogens, or pass `--prop-missing zero` deliberately.

The topomer transformation drops most of the molecules. The core is fused to the rest of them. Look at the cuttable fraction in the log and give a core joined to its R-groups by acyclic single bonds.

A decoy set gives suspiciously good enrichment. Check that the decoys are property-matched and posed in the training frame; cross-target decoys are separable by almost anything.

Co-fields are unavailable. A nonlinear kernel layer or the autocorrelation kernel source is selected, and neither is linear in the field, so the co-field does not exist. Use the *Shapley* tab instead.

It is slow. Kernel cost grows with the square of the total atom count, so the expensive sets are those with many large molecules. Cut the α grid first, since each α needs its own kernel, and remember that the count kernel is 10 to 34 times faster than the Gaussian route for a baseline.

Appendix A — Field parameterisation

Two parameterisations are available (*Increments* in the model panel, `--param`):

Field	`rdkit`, the default	`tripos`, the original CMF scripts
electrostatic	partial charge: Gasteiger, MMFF94, or read from the file	the same
steric	Tripos vdW well depth and width α/R^2_vdW, parameters by element	parameters by Sybyl atom type, Tripos 5.2 table
hydrophobic	Wildman–Crippen atomic logP contribution	in-house one-atom fragment increments by Sybyl type and attached hydrogens
H-bond acceptor / donor	approximate Abraham-type increments, donor weight on the polar hydrogen	in-house increments from Abraham's A and B constants, donor weight on the heavy atom
indicator field	1 for atoms of the given type	the same
`prop:<name>`	the value of a per-atom property carried by the file	the same

`tripos-legacy` reproduces the hydrogen-counting bug of the original R scripts, which is needed to match the published numbers exactly and should not be used for new work.

Appendix B — Files and tags the program understands

Tag or convention	Meaning
any numeric SD property	candidate activity; the loader guesses and the dialog confirms
`subset`, `set`, `split`	the train/test split, as `train`/`test` or 1/2
`atom.dprop.<name>`	a per-atom double property, written and read by RDKit
`atom.prop.sybyl_type`	per-atom Sybyl types, used automatically for indicator fields
a molecule property with one number per atom	the same per-atom property in quantum-chemistry dump form
`.cmf`	a saved model (joblib): field set, kernels, coefficients, training molecules, settings
Gaussian cube	the export format for co-fields, molecular fields, Shapley fields and difference maps
`REMARK 400` in a pocket PDB	the provenance of a trimmed binding site

Appendix C — Glossary

The terms this manual uses in a particular sense, with the chapter that treats each one.

term	what it means here
alignment, or **frame**	the common coordinate system the molecules are placed in before their fields are compared. The overlap kernel needs one and is only as good as it is; the count, autocorrelation, surface and periodic kernels do not. *Preparing the data*; *Working without an alignment*
applicability domain	how close a molecule outside the training set is to it, measured in the kernel the model itself uses, with a threshold taken from the training set's own spread. Reported per molecule as the **AD** column and carried into the exported predictions. *Where does the model apply?*
attenuation, α	the width parameter of a field's Gaussians: α large makes a field that falls away within an atom's own volume, α small one that reaches across the molecule. The single most consequential setting after the alignment. *1. The fields*; *Attenuation, regularisation and the grids*
autocorrelation kernel	geometry without a frame: the smoothed distribution of field-weight products over interatomic distance, invariant to rotation and translation by construction. *The rotation-invariant autocorrelation kernel*
BEDROC, RIE	screening figures that weight the top of a ranked list, where AUC saturates. *Decoys and the figures that matter*
Bemis–Murcko framework	the ring systems of a molecule together with the bonds that join them. `cmf diagnose` measures the common framework to ask whether one frame could exist at all. *Before you model*
block weights	one coefficient per block of channels — a block being one pair of field types — applied before the invariant kernels are normalised, so a weight of zero removes a block exactly. *Mixing coefficients over the channel blocks*
CHELPG, **Hirshfeld**, **Löwdin**, **Mulliken** charges	four ways of dividing a molecule's electron density among its atoms, all of which ORCA can report and any of which can become the electrostatic field's weights. They disagree, sometimes substantially, which is why the charge model is part of what must match between fitting and predicting. *Quantum-chemical properties with ORCA*
chemotype	a group of molecules sharing a scaffold. Splitting a benchmark by chemotype rather than at random is the test of whether a model carries beyond the series it was fitted on. *The example data sets*
co-field, $C_f(\mathbf{r})$	the model as a function of space: the training molecules' fields summed with the fitted coefficients, one per field type. Every prediction is the overlap of a molecule's own field with it. *The same prediction, as a field*; *Co-fields*
co-pattern	what the invariant kernels offer in place of a co-field — the weight of every channel of the fitted linear model, named by its field pair and distance, ranked by `top_channels`. *Crystals*; *Molecules on a surface*
contributions	the decomposition of one prediction over the training molecules, and of a training molecule's share over its atoms. *Contributions*

term	what it means here
count kernel	the pose-free baseline: a kernel on atom-type counts, with no geometry at all. Worth running before any aligned model, because it sometimes wins. *The pose-free count kernel*
decoys	presumed-inactive molecules used **only** to evaluate a screening model, never to fit it; property-matched decoys are the honest kind. *Decoys and the figures that matter*
`cmf diagnose`	the command that runs the five baselines — aligned, counts, autocorrelation, topomer, descriptors — on a data set before any model is built, and says which road to take. *Before you model*
dual coefficients, **a**, and **bias** b	what fitting produces: one coefficient per training molecule and one constant, from which every later prediction is a weighted sum of kernel values. *5. Fitting: what the training set buys*
ECFP4	a circular fingerprint of the atoms' environments out to two bonds: the standard 2D similarity measure, and the baseline the field similarity of this program is compared against on screening tasks, where it does about as well. *Decoys and the figures that matter*
ETKDG	RDKit's distance-geometry conformer generator, which builds the 3D structure whenever the input is 2D and supplies the starting conformations for the conformer ensembles, the flexible alignment and the topomer rebuilds. Its poses are arbitrary until they are aligned. *What the program reads*; *Conformations: the fourth dimension*
external cross-validation, q^2_{ex}	the whole model-building procedure, hyper-parameter search included, repeated inside each outer fold; the only figure that is not optimistic about the search. *Validation*
field, $X_f(\mathbf{r})$	a continuous function of space: one Gaussian per atom, weighted by that atom's charge, bulk, lipophilicity, H-bond strength, indicator value or file-borne property. The object CMF models. *1. The fields*
field scaling, c_f	the factor that puts a normalised field kernel's mean diagonal at 1 over the training set, so the fields enter the mixture on equal terms. Fixed on the training set and stored. *3. Normalisation and scaling*
field set	the list of field types a model uses, with their attenuations, charge model and parameterisation. It must match exactly between fitting and predicting. *What has to match between fitting and predicting*
flexible alignment	alignment that turns the rotatable bonds as well as moving the molecule, with the force field as a prior on the conformation and several generated starting conformations. The best pose method measured here: 0.86 Å median against the crystal poses. *Judging an alignment, and aligning flexibly*
Fukui index, f^+, f^-, f^0	a per-atom measure of where a molecule gives or takes electron density, computed from the N+1 and N−1 states by ORCA and used as a `prop:` field. The **dual descriptor** is $f^+ - f^-$. *Quantum-chemical properties with ORCA*
h, the **mixing coefficient** h_f	the weight of one field's normalised kernel in the joint kernel. Set by hand, optimised, or derived by kernel alignment. *4. Mixing the fields*
increments	the table of per-atom parameters a field is built from — the `rdkit` default or the `tripos` tables of the original CMF scripts. Not a setting of the method but of the chemistry it reads. *Appendix A*
indicator field	a field whose weight is 1 for atoms of a given type and 0 elsewhere, so that its co-field says where atoms of that type help. *Choosing the fields*

term	what it means here
joblib	the serialisation library a saved `.cmf` model is written with. A `.cmf` file is therefore a Python object graph, not an interchange format: it is read by this program and not meant to be read by another. *Saving and reusing*
joint kernel, $K(A,B)$	the sum of the normalised, scaled field kernels with their mixing coefficients: the single similarity the learner sees. *4. Mixing the fields*
Kabsch superposition	the least-squares rotation and translation that best superimposes two sets of matched points. It is the second half of an MCS alignment -- the substructure search finds the correspondence, this places the molecule -- and the pose benchmark measures how far the result lands from the crystal pose. *Alignment*; *Judging an alignment, and aligning flexibly*
kernel layer	a nonlinear function applied to the joint kernel — Gaussian or arc-cosine — after which there is no co-field, because the prediction is no longer linear in the fields. *Nonlinear kernel layers and multiple kernel learning*
MCS alignment	superimposing molecules on their maximum common substructure. The conventional frame, and the baseline every pose method here is measured against. *Alignment*
MKL	multiple kernel learning: deriving the mixing coefficients from the kernels and the activities by centred kernel alignment rather than setting them. *Nonlinear kernel layers and multiple kernel learning*
nested cross-validation	choosing the hyper-parameters inside each training fold and scoring outside it, so that the reported figure has not seen the molecules the choice was made for. *Validation*
normalisation	dividing a field kernel by the geometric mean of the two self-kernels, so every molecule has unit self-similarity and a large molecule cannot look important merely for being large. *3. Normalisation and scaling*
one-class model	a model fitted on actives alone, which asks whether a molecule resembles them rather than how active it is. *One-class models: screening with actives alone*
overlap integral, $K_f(A,B)$	the integral over all space of the product of two molecules' fields, which for Gaussians is a closed-form double sum over atom pairs. The kernel CMF is built on. *2. The overlap integral*
periodic kernel	the autocorrelation descriptor for a crystal: the pair histogram of one cell with its periodic images, per atom, with the uniform background removed and the cutoff tapered. *Crystals: the periodic autocorrelation kernel*
pharmacophore	a set of typed spheres extracted from the lobes of a co-field, which is the model's own statement of what a molecule should carry and where. *A pharmacophore out of the co-field*
probe field	a GRID-like interaction energy with a chosen probe, expanded in Gaussians so that it keeps the closed-form kernel and the co-field. *GRID-like probe fields*
`prop:` **field**	a field whose per-atom weights are read from the file or computed by ORCA rather than from a parameter table. *Charges and per-atom properties*

term	what it means here
q^2, R^2_{pred}, RMSE	the cross-validated, test-set and fit figures of merit. q^2 is not a test: it is what the search optimises, and a model can be chosen for it. *Validation*; *Statistics reported*
selectivity co-field	the difference of two endpoints' co-fields in a multi-task model, which draws where a molecule should differ to be selective rather than merely potent. *Two endpoints at once*
Shapley value	the share of a prediction attributed to each field or atom by averaging over all subsets, with the pairwise interaction indices beside it. *Shapley values*
solvent exposure weighting	multiplying each atomic weight by its relative solvent accessibility. Measured here and found not to help: ligands of this size have no interior. *Weighting the fields by solvent exposure*
surface kernel	the in-plane invariant kernel for molecules on a surface: distances in the plane, resolved by the heights of both atoms, invariant to lateral position and rotation about the normal and to nothing else. *Molecules on a surface*
topomer field	a field built from substituents rebuilt in a pose that is a deterministic function of their topology, so the same R-group always lands in the same place and no alignment is needed. *Topomer fields*
y-randomisation	refitting the whole procedure on shuffled activities many times; if the shuffled models score as well as the real one, the real one is chance. *Is it chance? y-randomisation*

Appendix D — Bibliography

References are grouped by what they are to this program: the papers that define the method it implements, those that apply it to one-class classification and virtual screening, and the work it borrows machinery or benchmark data from. Journal names follow the ACS abbreviations; a DOI is given wherever the work has one.

The method

[1] Zhokhova, N. I.; Baskin, I. I.; Bakhronov, D. K.; Palyulin, V. A.; Zefirov, N. S. Method of continuous molecular fields in the search for quantitative structure–activity relationships. *Dokl. Chem.* **2009**, *429*, 273–276. doi:10.1134/S0012500809110056

[2] Baskin, I. I.; Zhokhova, N. I. The continuous molecular fields approach to building 3D-QSAR models. *J. Comput.-Aided Mol. Des.* **2013**, *27*, 427–442. doi:10.1007/s10822-013-9656-4

[3] Baskin, I. I.; Zhokhova, N. I. Continuous molecular fields approach applied to structure–activity modeling. In *Application of Computational Techniques in Pharmacy and Medicine*; Gorb, L., Kuz'min, V., Muratov, E., Eds.; Challenges and Advances in Computational Chemistry and Physics, Vol. 17; Springer: Dordrecht, 2014; pp 433–459. doi:10.1007/978-94-017-9257-8_13

[4] Sitnikov, G. V.; Zhokhova, N. I.; Ustynyuk, Y. A.; Varnek, A.; Baskin, I. I. Continuous indicator fields: a novel universal type of molecular fields. *J. Comput.-Aided Mol. Des.* **2015**, *29*, 233–247. doi:10.1007/s10822-014-9818-z

[5] Baskin, I. I.; Zhokhova, N. I. Descriptors based on continuous indicator fields for 3D-QSAR studies. In *Proceedings of MOL2NET 2016, International Conference on Multidisciplinary Sciences*, 2nd ed.; MDPI: Basel, 2017. doi:10.3390/mol2net-02-03854

[6] Baskin, I. I.; Zhokhova, N. I. Continuous molecular fields and the concept of molecular co-fields in structure–activity studies. *Future Med. Chem.* **2019**, *11*, 2701–2713. doi:10.4155/fmc-2018-0360

One-class classification and virtual screening

[7] Karpov, P. V.; Baskin, I. I.; Zhokhova, N. I.; Zefirov, N. S. Method of continuous molecular fields in the one-class classification task. *Dokl. Chem.* **2011**, *440*, 263–265. doi:10.1134/S0012500811100016

[8] Karpov, P. V.; Baskin, I. I.; Zhokhova, N. I.; Nawrozkij, M. B.; Zefirov, A. N.; Yablokov, A. S.; Novakov, I. A.; Zefirov, N. S. One-class approach: models for virtual screening of non-nucleoside HIV-1 reverse transcriptase inhibitors based on the concept of continuous molecular fields. *Russ. Chem. Bull.* **2011**, *60*, 2418–2424. doi:10.1007/s11172-011-0372-8

Methods and data the program borrows

[9] Cramer, R. D.; Patterson, D. E.; Bunce, J. D. Comparative molecular field analysis (CoMFA). 1. Effect of shape on binding of steroids to carrier proteins. *J. Am. Chem. Soc.* **1988**, *110*, 5959–5967. doi:10.1021/ja00226a005

[10] Cramer, R. D. Topomer CoMFA: a design methodology for rapid lead optimization. *J. Med. Chem.* **2003**, *46*, 374–388. doi:10.1021/jm020194o

[11] Bemis, G. W.; Murcko, M. A. The properties of known drugs. 1. Molecular frameworks. *J. Med. Chem.* **1996**, *39*, 2887–2893. doi:10.1021/jm9602928

[12] Hopfinger, A. J.; Wang, S.; Tokarski, J. S.; Jin, B.; Albuquerque, M.; Madhav, P. J.; Duraiswami, C. Construction of 3D-QSAR models using the 4D-QSAR analysis formalism. *J. Am. Chem. Soc.* **1997**, *119*, 10509–10524. doi:10.1021/ja9718937

[13] Bishop, C. M.; Svensén, M.; Williams, C. K. I. GTM: the generative topographic mapping. *Neural Comput.* **1998**, *10*, 215–234. doi:10.1162/089976698300017953

[14] Cho, Y.; Saul, L. K. Kernel methods for deep learning. In *Advances in Neural Information Processing Systems 22 (NIPS 2009)*; Curran Associates: Red Hook, NY, 2009; pp 342–350.

[15] Rasmussen, C. E.; Williams, C. K. I. *Gaussian Processes for Machine Learning*; MIT Press: Cambridge, MA, 2006. doi:10.7551/mitpress/3206.001.0001

[16] Truchon, J.-F.; Bayly, C. I. Evaluating virtual screening methods: good and bad metrics for the "early recognition" problem. *J. Chem. Inf. Model.* **2007**, *47*, 488–508. doi:10.1021/ci600426e

[17] Sutherland, J. J.; O'Brien, L. A.; Weaver, D. F. A comparison of methods for modeling quantitative structure–activity relationships. *J. Med. Chem.* **2004**, *47*, 5541–5554. doi:10.1021/jm0497141

[18] Mittal, R. R.; McKinnon, R. A.; Sorich, M. J. Comparison data sets for benchmarking QSAR methodologies in lead optimization. *J. Chem. Inf. Model.* **2009**, *49*, 1810–1820. doi:10.1021/ci900117m